

Yasper

:

Yet Another Smart Process Editor

2.0

User Guide

:to be exact:

in the process of being updated for 2.0

What is Yasper?

Yasper is a tool for modeling and simulating stepwise processes. It is intended to be used for business process analysis and redesign, software requirements analysis and high-level software architecture design. Yasper intentionally limits itself to modeling the order dependencies in structures, stepwise processes, including issues of cooperation and resource contention, with some support for performance analysis; it is not intended for modeling other aspects of systems, such as data processing, organization structure, user interface design.

Yasper models rely on an exact definition of how processes can execute, even in the presence of concurrency, cooperation, and resource contention between processes. This makes more exact designs possible: errors can be spotted through *simulation*, symbolic execution of the model. This is achieved by defining processes in terms of *Petri nets*.

A Petri net describes a process by means of *places* and *transitions*. Transitions represent steps in the process: *actions* or *events*. Places represent conditions that hold before and/or after steps. Yasper extends Petri nets with notions such as *subprocesses*, *decision* blocks, *inhibitor* and *reset* connections, and *roles*. Yasper is particularly suited to the modeling of *workflows*: processes with designated start and end points. Unique to Yasper is the way in which it distinguishes between different *workflow cases* in a process. All modeling features are explained in chapter 1.

Since Yasper models have an exact execution semantics, simulation is a very powerful technique, both in model debugging and in the actual process analysis. Yasper supports two modes of simulation: *manual* mode, in which the user selects the subsequent steps to execute by clicking in the diagram (see chapter 3), and *automatic* mode, which randomizes the choice of steps and produces an aggregated report with relevant statistics.

Yasper was developed in collaboration between [TU Eindhoven](http://www.tue.nl) and [Deloitte](http://www.deloitte.com). For details of its status, availability and use, see the Yasper website:

<http://www.yasper.org/>

Yasper was written to provide modeling convenience, rather than new and unique new modeling features: the features it supports are well known, well analyzed, and support exists for them in other tools, in particular, in colored Petri net tools such as *ExSpecT*, *CPN Tools*, and *ReNeW*. These tools can additionally express many things that Yasper does not support. What Yasper aims to do is to make modeling and simulation easier to perform and to learn, especially for workflows. Yasper models can emulate techniques such as state machines, flowcharts, and activity diagrams.

How to read this User Guide

This user guide assumes you have decided to give Yasper a try, and attempts to provide you with everything you need know in order to use it. It does not assume any prior knowledge of the techniques it uses; however, where appropriate, some comparisons with similar tools and techniques will be made for the benefit of the reader who is already familiar with them. Depending on your background knowledge you may wish to skip certain parts of this guide.

Chapter 1 describes Yasper models. It details Yasper's modeling constructs, how they fit together, how they appear to the user, and what they mean in terms of process execution.

Chapter 2 describes the Yasper editor. It details the user interface functionality related to the specification of models. It also describes the means available to import or export Yasper models from/to other applications.

Chapter 3 describes manual execution mode. It details the user interface functionality specific to manual execution and describes the relationship between this mode and edit mode.

Chapter 4 describes automatic simulation mode. It details the user interface functionality related to setting simulation parameters and executing a simulation run, and defines the meaning of the statistics reported.

Table of contents

What is Yasper?	2
How to read this User Guide	3
Table of contents	4
1 Process models in Yasper	6
1.1 Basic nets	7
1.1.1 Basic net elements	7
1.1.2 Process execution in basic nets	8
1.1.3 Modeling with basic nets	9
1.1.4 Transitions and places must alternate	13
1.2 Choice elements	14
1.3 Roles	16
1.4 Hierarchical models	18
1.5 Special arc types	20
1.5.1 Combined arcs	20
1.5.2 Inhibitor arcs	21
1.5.3 Reset arcs	21
1.6 Extensions for automatic simulation	23
1.6.1 Arc weights	23
1.6.2 Processing time	23
1.6.3 Processing cost	24
1.6.4 Token case	25
1.7 Token color	32
1.8 Data stores	33
2 The Yasper editor	34
2.1 Starting with an empty model	35
2.2 Selecting and deselecting elements	37
2.3 Moving elements	38
2.4 Merging elements	41
2.5 Deleting, copying and pasting elements	42
2.6 Connecting nodes	43
2.7 Extending the node frontier	44
2.8 Arc support points	45
2.9 Changing the type of an element	46
2.10 Specifying additional properties	47
2.10.1 Transition properties	47
2.10.2 XOR properties	49
2.10.3 Place properties	50
2.10.4 Arc properties	51
2.11 Working with subnets	53
2.11.1 Creating, navigating, and removing subnets	53
2.11.2 Connecting across subnet boundaries	54
2.11.3 Working with the tree view	55
2.12 Working with roles	57
2.12.1 Defining roles	57
2.12.2 Assigning roles to transitions	57
2.13 Toolbar operations	60
2.14 <i>Edit</i> menu operations	61

2.14.1	Undoing changes	61
2.14.2	Copying and pasting	61
2.14.3	Aligning elements	62
2.14.4	Making elements evenly spaced	63
2.14.5	Snapping to the grid	64
2.15	<i>File</i> menu operations	65
2.15.1	Loading and saving Yasper models	65
2.15.2	Importing BPMN models	67
2.15.3	Exporting models to Microsoft Visio	67
2.15.4	Exporting diagram images	68
2.15.5	Printing diagrams	69
2.16	<i>View</i> menu options	71
2.17	Zooming in and out	73
2.18	<i>Options</i> menu options	74
3	Manual simulation	75
3.1	What is simulation?	76
3.2	Preparing for simulation	77
3.3	Starting manual simulation	81
3.4	Executing a manual simulation	83
3.5	Reviewing a simulation	87
3.6	Manual simulation with roles	88
4	Automatic simulation	90
4.1	How automatic simulation works	91
4.2	Preparing to run a simulation	92
4.3	Not being able to run a simulation	93
4.4	Starting, halting and restarting a simulation	94
4.5	The simulation report	96
4.6	Switching views and modes	98
5	Index	99

1 Process models in Yasper

This chapter describes the process modeling features supported by Yasper.

If you just want a description of how to operate the Yasper tool, skip to chapter 2 (the editor), 3 (manual process execution) or 4 (automated process execution).

Yasper's process models are *Petri nets*. Like most Petri net tools, it extends basic nets, defined by C.A. Petri in 1962, with additional constructs.

Section 1.1 describes basic Petri nets: networks of *transitions* (process steps or events) connected through *places* (conditions that must be fulfilled before or after a process step). Examples of typical modeling patterns are given.

Section 1.2 adds *choice*, a construct known from flowcharts and other techniques.

Section 1.3 adds *roles*, an alternative to using places for modeling resources.

Section 1.4 adds hierarchical modeling with the subnet construct.

Section 1.5 adds special types of connections: combined arcs, inhibitor arcs and reset arcs.

Section 1.6 describes arc weights.

Section 1.6.2 and 1.6.3 add processing time and processing cost.

Section 1.6.4 adds *token case*. With this feature, process execution can distinguish between different cases being processed at the same time. This makes it easier to model and simulate workflow-like processes.

Sections 1.7 and 1.8 describe Yasper's rudimentary support for *token color* and *stores*.

1.1 Basic nets

This section defines Yasper's support for basic Petri nets – known in the literature as place/transition nets, or P/T nets. For a rigid introduction to Petri nets and their properties, the reader is referred to the many introductory books and articles that have appeared on them.

1.1.1 Basic net elements

Models express the flow structure of processes. The simplest kind of model only uses three constructs and is called a *basic net*.

1.1.1.1 Transitions and places

The most important modeling unit represents a process step. It appears in the diagram as a green square.



fuel the car

Process flow structure means that not every step can occur at every moment: they are subject to certain conditions. In Yasper, these conditions are made explicit.




car out of
fuel

The squares represent actions, processes or events; generally speaking, *transitions* between states of affairs in the world.

The circles represent conditions, and together, the possible states of affairs as far as they are relevant to the process flow of the process being described. They are known as *places*.

1.1.1.2 Arcs and nets


Transitions and places are connected to form a *net*; for example,



The arrows are known as *arcs*.

1.1.1.3 Tokens and markings

A state of affairs is represented by marking the conditions that hold:



The dots are known as *tokens*; a distribution of tokens over places is a *marking*.



1.1.2 Process execution in basic nets

Process execution is defined by a single rule: a transition can occur whenever all places that have arcs to it contain a token; its occurrence will remove a token from each of them, and put a token in each of the places to which it has an arc.

For example, in the above diagram, *fuel the car* can only occur when *car out of fuel* is marked with a token, and when it occurs, the token moves into *car full of fuel*.


As soon as the token arrives there, *car is driven around* becomes possible; when it occurs, the token moves back into *car out of fuel*. So we see the model is not very realistic: a car cannot be driven without running out of fuel, and it cannot be fueled before it has run out. But the point of the example is that process execution is completely defined by this one transition execution rule.

Petri net modeling tools, such as Yasper, allow transitions to be executed in the diagram, either stepwise (see chapter 3) or automated (see chapter 4). Because the user interactively moves token around in the diagram, this is known as *playing the token game*.



In this model, *person eats food* can only happen when a person is hungry and food is available; and when it does, the person is no longer hungry, but satisfied, while the food is no longer ready, but gone.

It is legal for multiple tokens to be in the same place:



In this example, the tokens represent cars: three still out of fuel, one ready to drive. Generally speaking, in models using this technique, tokens usually describe objects, while places describe states these objects can be in. A whole marking describes the overall state of the system.

Similarly, any number of persons and food portions can be introduced into the other example.

Note that execution is nondeterministic: it is precisely determined what happens when a transition executes, but when multiple such executions are possible at the same time, it is not determined which one will happen. More on this in section 1.6.4.4.

1.1.3 Modeling with basic nets


This section presents some examples to illustrate typical modeling patterns.

The intention is to provide some basic notions and examples that later sections will build on to introduce additional Yasper features.

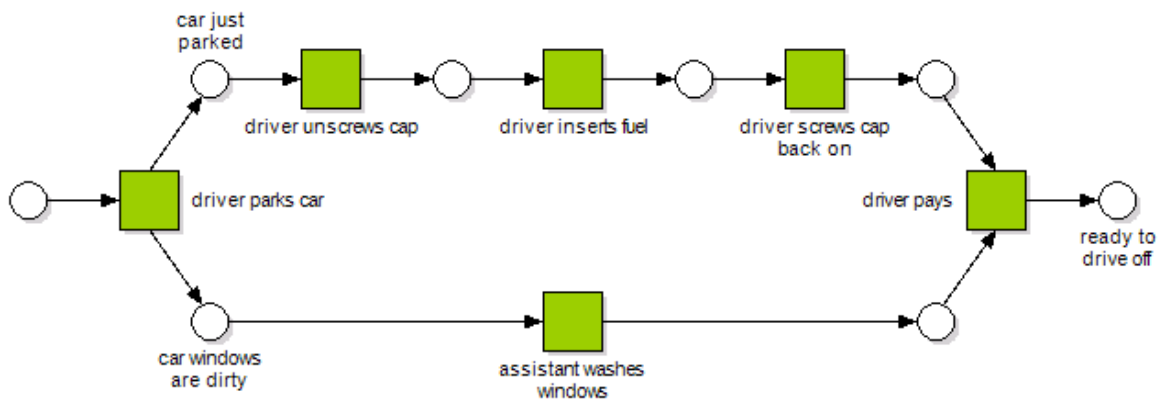
Full tutorials on the use of Petri nets can be found elsewhere.

1.1.3.1 Three useful types of branching in basic nets

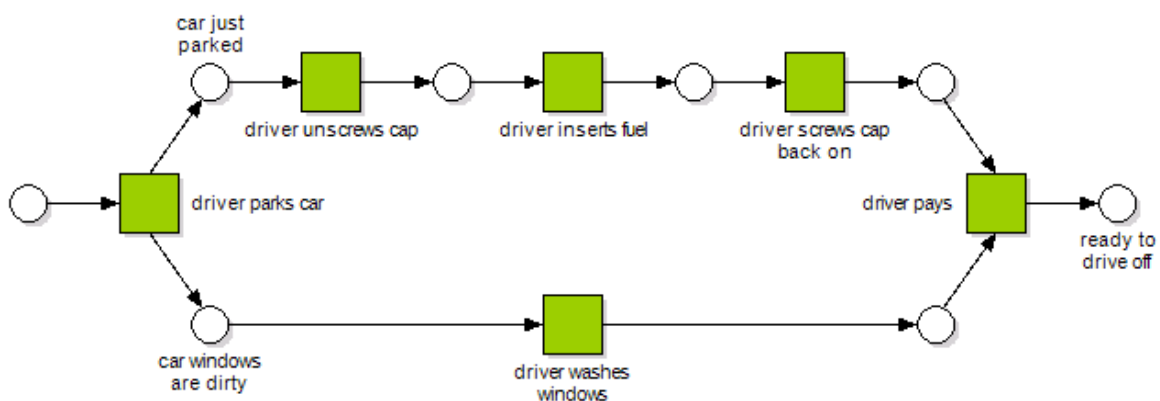
Alternative paths can be described by branching on places:



Parallel or concurrent execution, on the other hand, is described by branching on transitions:

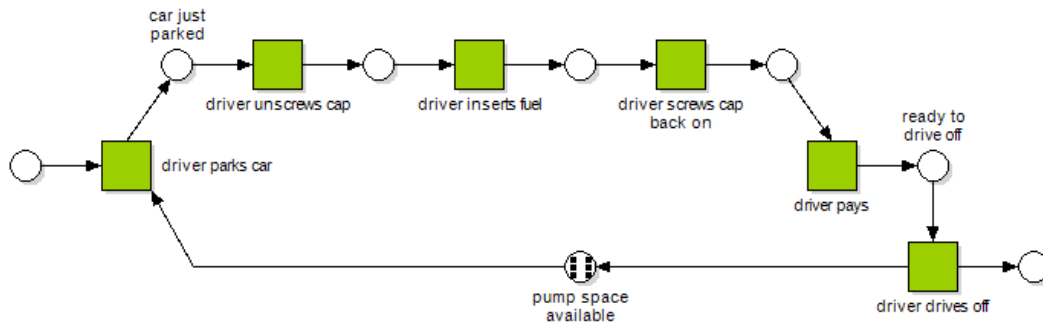


This introduces *nondeterminism*: the order in which the upper and lower branch execute is undefined. For example, the assistant may finish before the driver has unscrewed the cap, or may start washing while the driver is inserting fuel and only finish after the driver is ready, or any other combination. In fact, any situation in which the order of steps is undefined must be modeled with parallel branches:



In this example, there is no implication that things are done in parallel; it just models the situation in which the windows can be washed at any time while the other three steps must be performed in the order given.

Branching on transitions can also be used to describe resource contention:



The execution rule given above determines that whenever six cars have parked and remain somewhere in the process, *driver parks car* will not be possible until one of them drives off, thereby releasing the space it occupied. All incoming cars would be kept waiting, according to this model.

1.1.3.2 Well-nested nets and proper completion

As the examples show, these few basic constructs and the one execution rule give us a very simple, yet very powerful technique to describe process flow logic. Nice aspects are that the flow logic is fully defined by the execution rule, and that it is fully displayed in the diagram.

The three examples have some things in common:

- all branch points form pairs of splits with matching joins of the same type; if a net has this property, it is called *well-nested*,¹
- they describe *workflows*, processes in which individual cases (represented by tokens) flow from a fixed starting point to a fixed end point; if such a net can handle such business cases indefinitely and concurrently without ever locking up or amassing tokens anywhere, we say it *completes properly*.²

An exact definition of proper completion is given in section 1.6.4.5.

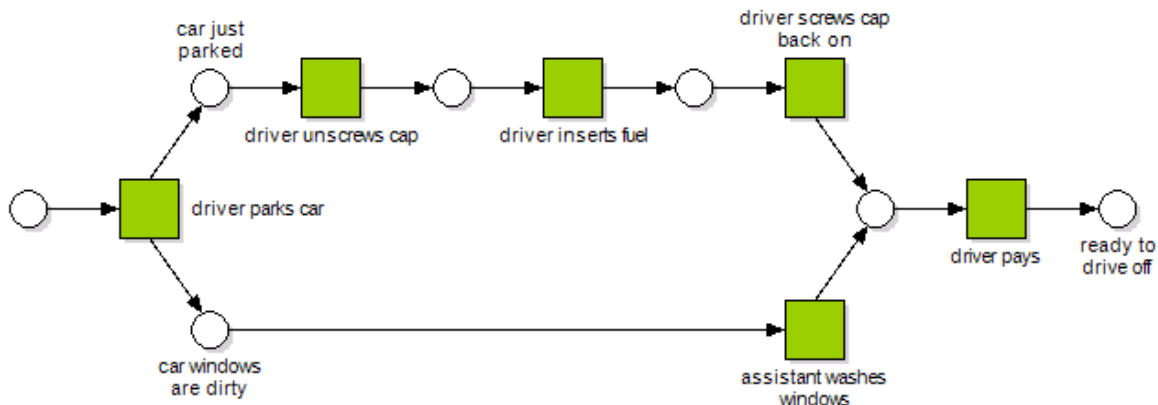
¹ Somewhat looser criteria on the structure of nets have been studied, such as *well-handled nets* (see: *Workflow Verification: Finding Control-Flow Errors Using Petri-Net-Based Techniques* (Van der Aalst, 1998) . *Well-structured nets* are yet another notion.

² Proper completion generalizes the *soundness* property defined in *Workflow management: Models, Methods, and Systems*, (Wil van der Aalst and Kees van Hee, 2002), which requires the net to be empty initially and eventually.

In the practice of modeling, proper completion is an important correctness criterion.

A net with a clear begin and end point that lacks this property may well point out a real problem with the process in question, but more often, is an incorrect model of the process. It turns out to be very easy for human modelers – including the authors of Yasper - to make mistakes even in simple models, and checking for proper completion turns out to reveal most such errors.

Well-nested nets always complete properly, so it is a good design principle to stick to neatly bracketed split-join pairs whenever possible. Doing so makes problems easy to spot, as can be seen in this beginners' error:



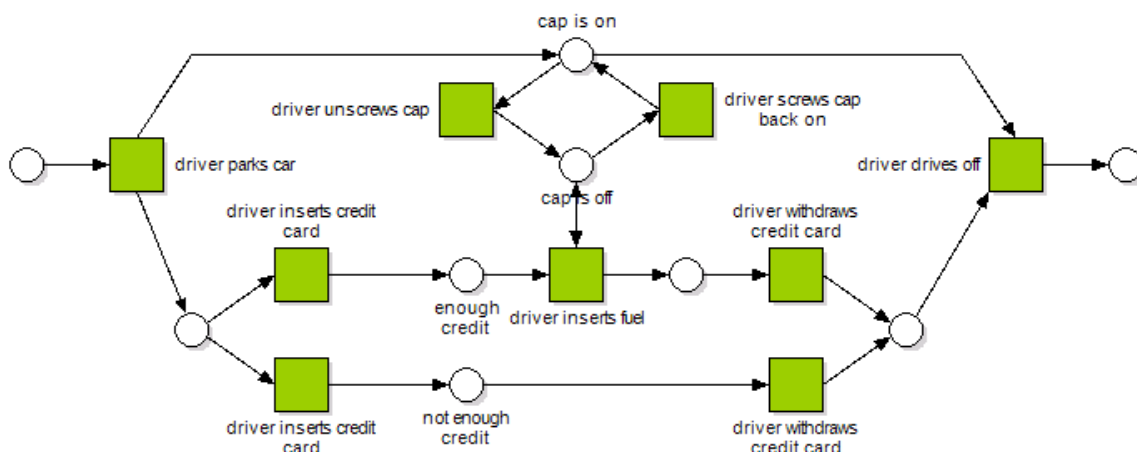
Every driver parking a car will pay twice and end up with two cars!

The cause is obvious: a transition split is mismatched with a place join.

Simulation will make this clear – see section 3.2 for details.

However, in well-nested nets, the only form of parallelism is completely hierarchical: a process can start multiple concurrent processes and wait for all of them to finish, but communication or synchronization between different subprocesses is impossible. In practice, a large class of processes can only be modeled with nets that properly complete, but are not well-nested.

This situation can arise as soon as the constraints on the order of process steps are anything other than a choice between fixed paths, even when they are quite simple:



At this level of complexity, verifying proper completion already becomes a problem to the human eye. Proper completion can in fact be verified automatically in many cases. However, Yasper doesn't yet offer such verification, because it has a much more insightful way to check for improper completion: using manual or automatic simulation, the user sees the tokens flowing through the net and spots deadlocks or bottlenecks where they arise. Simulation turns out to be an extremely convenient tool in designing correct non-well-nested nets, and appears to catch nearly all of the modeling errors that arise in practice.

1.1.4 Transitions and places must alternate

In Petri nets, arcs always run between transitions and places: it is illegal to connect two transitions or two places directly. Such arcs could be introduced as shortcuts, on which the intermediate place or transition has been omitted.

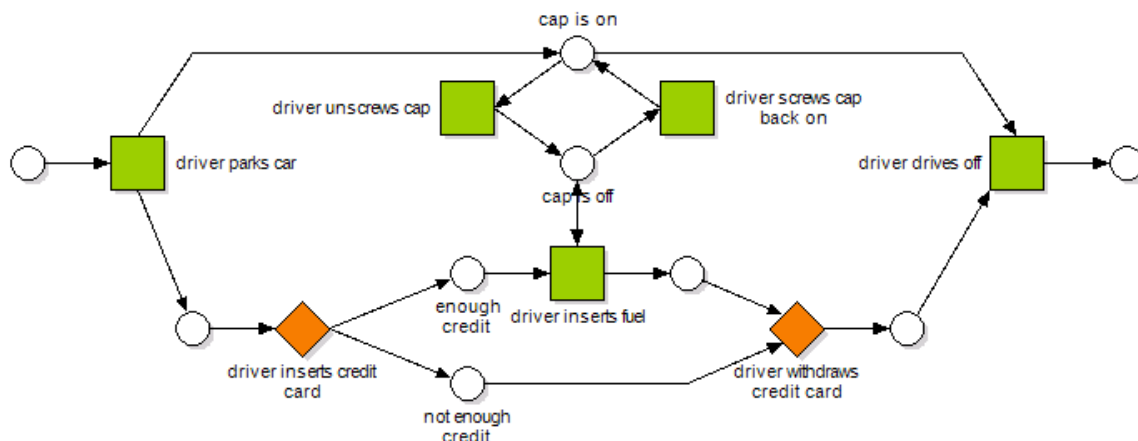
Although we have experimented with this feature, the present Yasper release does not support such shortcuts. It is tidier to always make all transitions and places explicit, since branching can happen on both.

In practice we often see nets in which all branching happens on places: such models express alternative paths, but no parallelism or resource contention. A net with this property is called a *state machine*. For them, it makes sense to always omit the transitions and connect the places directly; this is in fact how state machines are usually drawn. Historically, state machines pre-date Petri nets, and they are still used as a basis for popular specification techniques, such as UML statecharts.

Many process modeling techniques exist today, often based on state machines (e.g. UML statecharts) or flowcharts (see the next section). Some of these techniques can easily be described with Petri nets: e.g., UML 1 activity diagrams and the XLANG or BPEL process models used in Microsoft BizTalk are effectively well-nested Petri nets without initial tokens; as we have seen, this implies that they cannot directly express interprocess communication or resource contention.

1.2 Choice elements

There are two places in which the preceding example looks quite awkward: two of the process steps appear twice. The reason is that the steps in question are subject to varying conditions: for *driver inserts credit card*, its results, and for *driver withdraws credit card*, its prerequisites. The process execution rule does not allow either, and duplication is the only way to resolve this. Clearly it is more natural to just write:




The new symbol is called a *choice* or *XOR* element, a construct used in *flowcharts*, a technique developed for computer programming by Goldstine and Von Neumann in 1946-1947,³ but well established in engineering much earlier;⁴ and in modern derivatives of the flowchart such as UML activity diagrams and BPMN.

The process execution rule is extended as follows: a choice consumes a token from one of its input places, if any, and produces a token in one of its output places, if any.


Again, it is best practice to use XORs in pairs, so as to keep models well-nested, but this is not always possible. For instance, they often appear in loops:

³ See pp. 266-267 in Herman.H. Goldstine, *The Computer from Pascal to Von Neumann*, Princeton University Press, 1972, ISBN 0-691-08104-2

⁴ See Gilbreth F.B. and Gilbreth L.M., (1921) *Process Charts - First Steps in Finding the One Best Way to do Work*. Presented at the Annual Meeting of The American Society of Mechanical Engineers, New York, USA.



An XOR can be defined as a place surrounded by non-branching transitions:



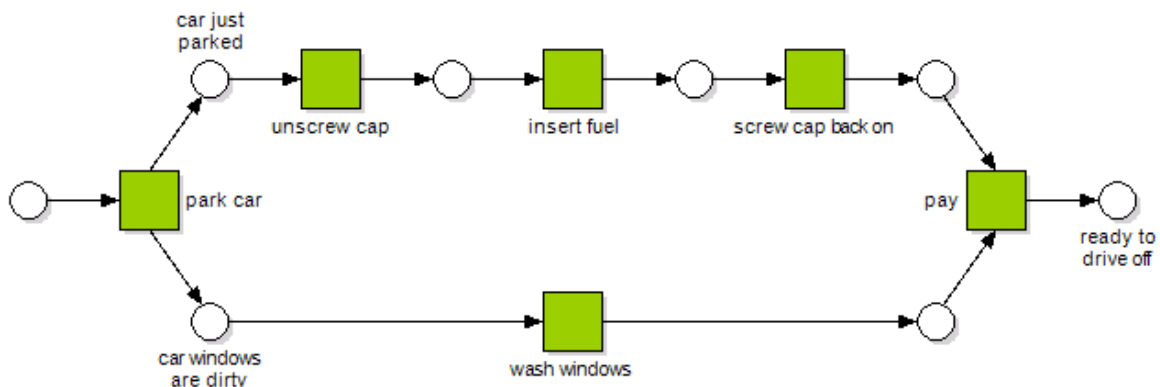
At present, Yasper doesn't quite treat XORs in that way, since it allows roles (cf. section 1.3) and output weights (cf. section 1.6.1) to be associated with them, and treats their execution atomically in manual simulation (cf. section 3.4).

1.3 Roles

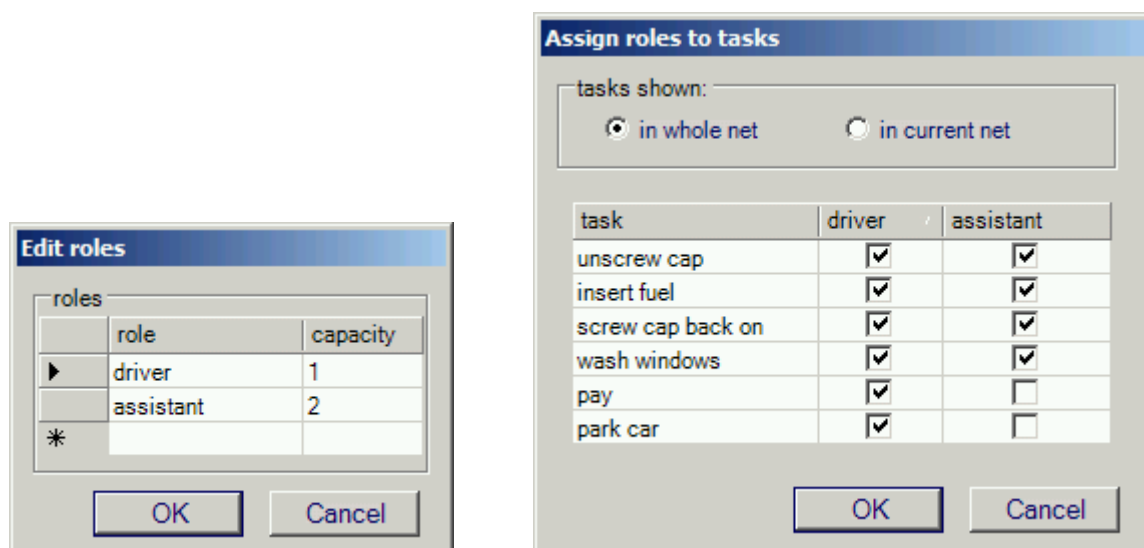
In business processes, production processes, or computer algorithms, efficient resource allocation is often a major driving force behind process design. Yasper offers a special facility for this purpose: the execution of a transition can be assigned to one or more *roles*.

Consider, for instance, the parallel execution example in section 1.1.3.

We may wish to express that in general there are two types of agents available to perform the tasks involved: drivers and assistants. In Yasper we can define *driver* and *assistant* as roles, and tag individual transitions with the roles that can execute them. The general model would look like this:



Roles do not appear in the diagram, but are defined separately (see section 2.12):

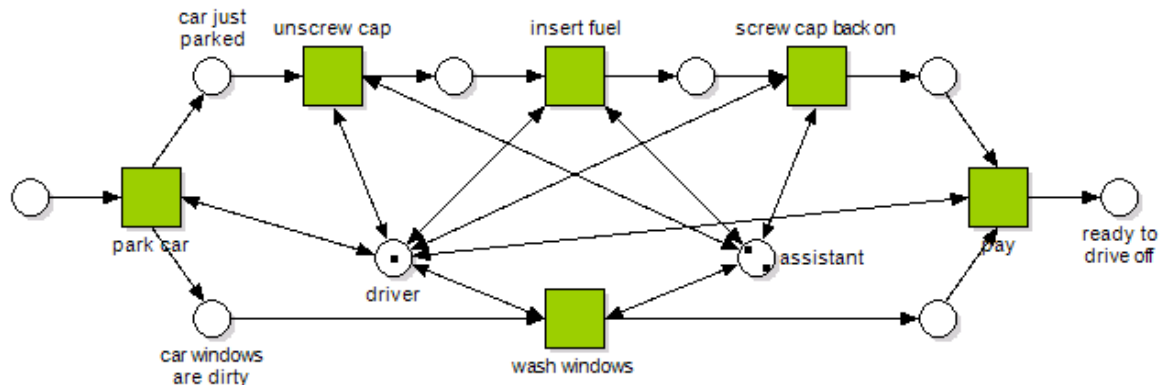


Note that roles have capacities: there is one driver and there are two assistants.

Process execution is influenced by roles in the expected way: a transition with assigned role(s) can only execute when an instance of a role is available to perform it, and while it executes, will keep that instance from doing something else.

For example, the windows can only be washed if a driver or an assistant is available, and no more than three cars can have their windows washed at the same time.

Although resources can also be modeled with places, a role is not equivalent to a place. Consider the following example:




The capacities and assignments to transitions are the same, but execution is different: in this model, *wash windows* occupies both a driver and an assistant, while in the previous model, one of either suffices.⁵

Whether to model resources with places or with roles depends on the purpose of the Yasper model. Manual and automatic simulation have special support for roles; for example, role utilization is aggregated in the automatic simulation report. See sections 3.6 and 4.5 for details.


⁵ Note that neither example is realistic: in practice, each car comes with its own driver, so drivers should not be modeled as resources at all.

1.4 Hierarchical models

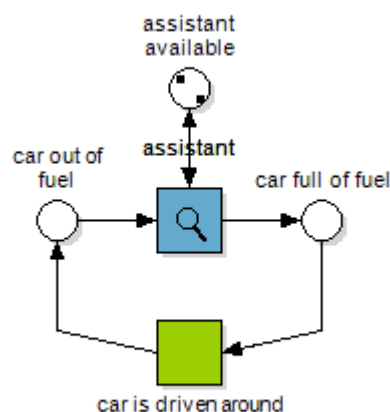
Models can grow big rather quickly. It is possible to spread the contents of a net over multiple pages by means of the *subnet* construct. In the surrounding net, a subnet is like a transition, in that it must be connected to places:




The contents of the subnet are those of a normal net, except that they may contain *pins*. A pin represents an arc connecting the subnet in the surrounding net. For example, the content of the subnet above may look like this:



The left and right pins correspond to the incoming and outgoing arc, respectively. Ambiguities can be clarified by naming:





Note that two pins can refer to the same place, and that pins cannot hold tokens: the tokens shown here are those of the place they refer to, *assistant available*.

Subnets can occur within subnets: in that case, a pin can refer to a pin in the surrounding net.

Execution with subnets is straightforward: all pins are identified with the places they ultimately refer to. Hence, subnets are a purely visual device to allow large diagrams to be split across multiple pages.


1.5 Special arc types

There are three special types of arcs: *combined arcs*, *inhibitor arcs*, and *reset arcs*.


1.5.1 Combined arcs

Most definitions of Petri nets allow multiple arcs to run between the same transition and place. In Yasper, all input and output arcs between the same transition and place are bundled into a single arc.

The *biflow* arc represents a pair of input and output arcs; it already occurred in some of the previous examples. An arc can represent any number of input and output arcs; in general, both numbers appear on the arc:




This is rare in practice, but multiple inputs or multiple outputs are more common:



Here, *fuel the car* always brings 25 gallons of fuel from the station tank to the car tank.

A combined arc with one or more multiplicities is usually equivalent to multiple single arcs. In the present Yasper version this is not the case for arcs on XORs.



If the arc between p and x behaved like three separate arcs, an execution of x would always pick a single token from p or q, and put a single token into p or r. Instead, when x decides to consume

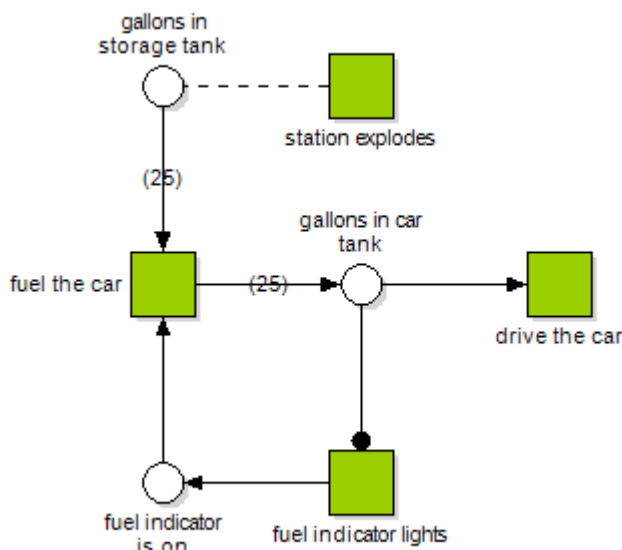
from p , it always takes two tokens from it. It does not guarantee to put a token back into p in that case; it may put it into r instead.

This behavior is odd and may change in the next Yasper version – see section 1.6.

The arcs we have seen thus far are called *flow arcs*, since they represent tokens flowing between transitions (or XORs) and places.

1.5.2 Inhibitor arcs

Let's extend the previous example a little:



The arc with the dot on the end is an *inhibitor arc*. It specifies a constraint on execution, namely, that the transition at its end can only execute when there are no tokens in the place at the other end. In the example, the fuel indicator lights when the gas tank is empty.


The inhibitor arc is very powerful⁶, but it takes skill to apply it correctly; in fact, there is a problem with the way it is used in this model – see the following section.

1.5.3 Reset arcs


A reset arc does not express a constraint, but instead specifies that when the transition at its end executes, it empties the place at the other end.⁷ In the example, after the fuel station explodes its storage tank will be empty.

⁶ Arbitrary computations on numbers of tokens can be expressed with inhibitor arcs: they make Petri nets Turing complete.

Yasper draws a reset arc as a dashed line, or, alternatively, as a 0-marked arrow:



The example as given requires some corrections:



In the original, once the fuel indicator can go on, it can continue to do so, producing an arbitrary number of tokens in *fuel indicator is on*. But we only have one fuel indicator, and while it is on it cannot switch on again.


Furthermore, in this new version, the station can only explode when its tank is nonempty. It also provides some initial fuel, without which nothing can ever happen.

⁷ *Reset arc* is the name used in the Petri net literature. It can be argued that a reset arc doesn't really perform a *reset*, since it clears the tokens from a place, even when that place wasn't empty initially.

1.6 Extensions for automatic simulation

1.6.1 Arc weights

During simulation, when a XOR split is executed, an output place must be chosen. By default, every place has equal likelihood to be picked.



In this example, the modeler may wish to indicate that when we look at a price, it has a 90% chance of being acceptable. To this end, the modeler can attach *arc weights* to the output arcs of a XOR. It is not possible to do the same for input weights.

Arc weights only take effect in automatic simulation; during manual simulation, the output place is picked by the user.

Arc weights are not displayed in the diagram, and must be edited by opening the property sheet of the XOR and setting the arc weights there. They serve a different purpose than arc multiplicities, discussed in the previous section. In the next version of Yasper we may simplify this situation, and use arc multiplicities on XORs to express arc weights.

1.6.2 Processing time

Two parameters can be set on each transition to specify the time its execution takes

- its mean processing time
- the mean deviation of its processing time

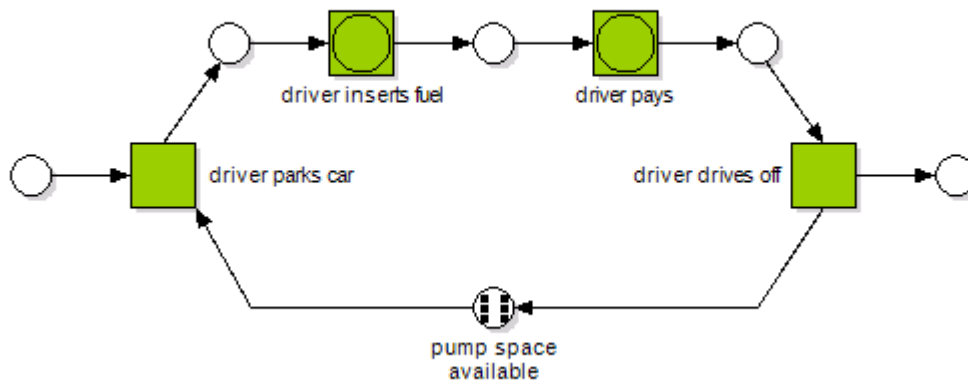
During automatic simulation, the actual processing time spent is determined by these two parameters. If the given deviation is 0, the processing time is always equal to the given mean. Otherwise, whenever the transition starts execution, a processing time is chosen randomly such that on average the chosen value will be the given mean, and the standard deviation from the mean will be the given deviation. This is achieved by using the so-called *gamma distribution*.

Note that processing time solely depends on these two parameters of the transition being executed; no other information can be used. E.g., to express that the same action can be executed by either a slow or a fast machine, place two transitions in parallel.


Timed transitions without inputs fire right at the start of a simulation; from then on, they time out to each subsequent firing, according to the time parameters.

Processing time specifications can strongly influence the paths a simulation will take. It is easy to construct nets that are perfectly viable in principle, yet completely lock up with certain choices of processing times. Often such a lockup turns out to reflect a real bottleneck in the process being modeled, and experimenting with processing times is a powerful tool in designing efficient process flows.

In the diagram, timed transitions (with *mean* > 0) are drawn with a circle; the exact parameter values are not shown.



This method of modeling time was chosen because it is intuitive, especially to users with past experience in flowcharts or activity diagrams. But the resulting “transitions” are no longer immediate, atomic state changes, as is customary for Petri nets. They are better thought of as a pair of transitions: one that starts the activity, and one that completes it:



Yasper’s execution engine also implements them in that way.

1.6.3 Processing cost

Processing cost is also configurable with two parameters per transition:

- its fixed (setup) cost **f**
- its cost per time unit **v**

The cost of a process execution is $f + v \cdot t$, where t is its processing time.

Processing cost cannot influence process execution at all – there is no provision in Yasper to make execution paths depend on cost. Cost calculations only appear in the automatic simulation report, where they provide indications of the expected overall cost of a particular process configuration.


Processing cost is not shown in the diagram.

1.6.4 Token case

The examples we have seen so far are *workflows*: the process takes tokens from a starting point to an end point. Such tokens represent cases being handled by the workflow. This is a very common type of model.

Yasper's automatic simulation feature (see chapter 4) was designed for workflows. It generates workflow cases, runs them through the model, and produces a report detailing how the average case was handled.

Automatic simulation requires that the start, continuation and end of workflow cases are clearly indicated in the model. This must be done with three specifically designed constructs: *emitters*, *case sensitive places*, and *collectors*, respectively. The smallest model on which automatic simulation will work has one of each:



1.6.4.1 Case sensitive places

A *cased* token belongs to a specific workflow case; an *uncased* token doesn't.



In Yasper, a place can be marked as *case sensitive*, which means it can only hold cased tokens. Such a place is shown in yellow. Any other place can only hold *uncased* tokens.


1.6.4.2 Emitters

An emitter is a transition without case sensitive inputs and with one or more case sensitive outputs. Every firing of an emitter establishes a new workflow case, and marks the tokens with it that are emitted to its case sensitive output(s). Transitions pass the token's case on through the net, as we will see. Therefore, a workflow case is defined by the emitter that generated it and the generation step at which it was generated.




1.6.4.3 Collectors

A collector is a transition with one or more case sensitive inputs and without sensitive outputs. Workflow cases, started at emitters, are expected to finish at collectors, although it is possible to create models in which this does not happen. Automatic simulation reports on workflow cases flowing from emitters to collectors (cf. section 4.1).



The simulator supports models with multiple emitters and/or multiple collectors. This can be used for simulating models in parallel.

An XOR can be an emitter or collector, so long as it is unambiguous. In other words, an XOR cannot have a mix of case sensitive and case insensitive inputs or outputs:




Here, it is undefined how *t* will execute. Yasper's execution engine happens to accept such nets and exhibits some behavior, but the situation is nonsensical and the model should be corrected.

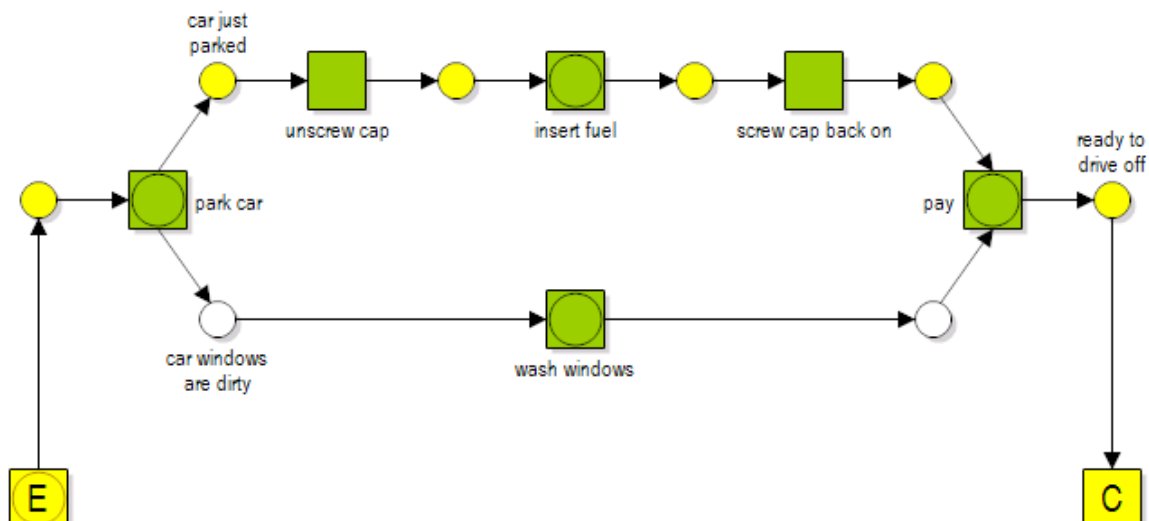
1.6.4.4 Case sensitive execution

Transition execution works on a case-by-case basis. To be exact, the execution rule is adapted as follows: all tokens consumed from case sensitive places must be of the same case, and that case will be passed on into any case sensitive output places.

The examples so far were workflows, so they are easily modified to generate automatic simulation reports.

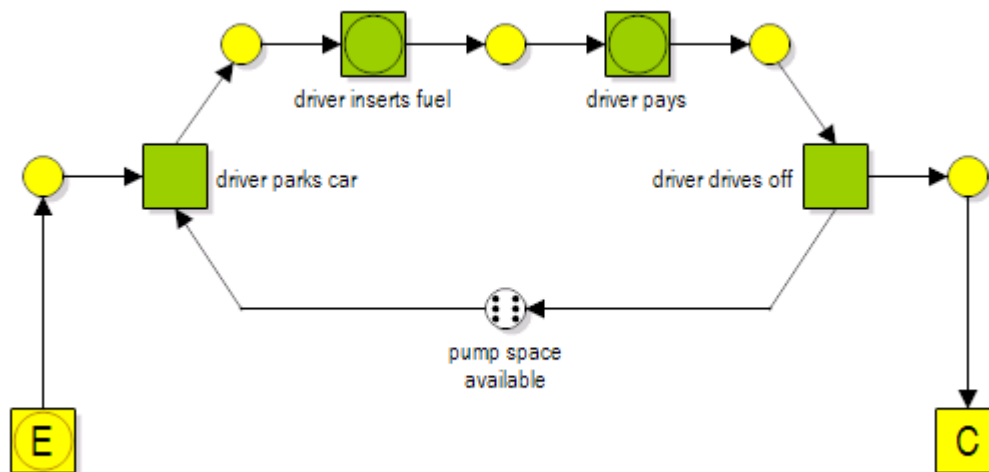


The case-by-case behavior of transitions becomes relevant as soon as multiple cases enter the model. In this example, *pay* is only possible when fueling and washing have both finished for the same car. Contrast this with the following:




Here, *pay* is possible whenever fueling has finished and washing has finished for some car. Clearly, not the intended semantics, and if *wash windows* takes a relatively long time, the difference will be clearly noticeable in the simulation report.

A common use of case insensitive places is to model resource contention.



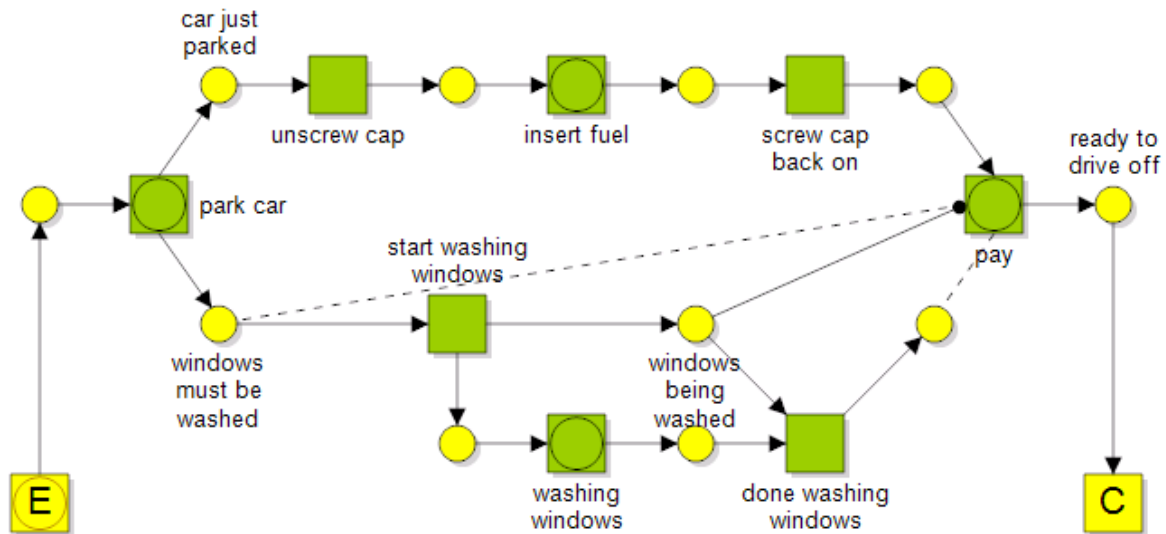
Here, *pump space available* must be case insensitive, since it describes how cases affect each other. If it was case sensitive, *driver parks car* could never execute.

Inhibitor and reset arcs also work case-by-case:



Here, we pay and drive off, whether or not the windows have been washed. Since the resets are by case, payment for one car doesn't affect whether the windows of other cars are washed.


However, *pay* also proceeds if the car was being washed, leaving a token behind in wash windows that will never be cleaned up: this model does not *complete properly* (as defined in section 1.1.3.2). Not only will this cause problems for the simulation report, we actually want to wait when the car is being washed.



Note that the inhibitor only suspends payment while the *same* car is being washed. To suspend it while *any* car is being washed, *windows being washed* must be case insensitive.

We now provide an exact definition of what it means for a transition or XOR to execute. As in the case of basic Petri nets, this definition completely defines process execution – there are no other constraints or possibilities.

In general, a transition may have any combination of cased and uncased places attached to it with any combination of arc types, and with any assignment of roles, time and cost.



Execution of a transition can start whenever the following conditions are met:

- if any roles are assigned to it, a role is available
- all case insensitive places connected to it with inhibitor arcs, if any, are empty
- every input place contains as many tokens as there are arcs from it to the transition (counting multiplicities), all having the same case on the case sensitive input places, such that

- no case sensitive places connected to it with inhibitor arcs contain any tokens with that case


When execution of a transition starts, the following happens, atomically:

- a set of tokens on the input places as just described is consumed (removed from their input places)
- if any roles are assigned to the transition, one is occupied (i.e. its availability, initially equal to its capacity, decreases by 1)
- all tokens are removed from case insensitive places connected to the transition with reset arcs
- if any cased tokens were consumed, all tokens with that case are removed from case sensitive places connected to the transition with reset arcs
- the processing time for this job is determined
- the job is started

Note that a transition without input places can still consume (the set of tokens consumed is empty).


When the processing time of a job has expired, and also, for inputless transitions, at the very start of execution, the following happens, atomically:

- if it occupied a role, it is released (i.e. its availability increases by 1)
- to every output place as many tokens are added as there are arcs to it from the transition (counting multiplicities), with the same case if the place is case sensitive, and without case otherwise, the case being equal to the case of the cased tokens consumed for this execution, if any, and a newly generated case, otherwise



For XORs the rules are identical, except that XORs consume from only one of the input places, if any, and produce to only one of the output places, if any. The XOR depicted here is doubly impossible, since it combines case sensitive and case insensitive inputs, and also combines case sensitive and case insensitive outputs.

The following (impossible) model illustrates why this is invalid:



Here, when t consumes from $p1$ it acts as an emitter, starting a new workflow case in $p2$, but when it consumes from $p3$, it will continue the case of the token consumed. This seems useless and rather confusing, so such XORs are forbidden.

Multiple transitions may simultaneously be able to start and/or finish jobs, and a transition may simultaneously be able to start and/or finish different jobs. In Yasper's manual and automatic simulation modes, finishing always takes precedence. Furthermore, whenever a transition can consume for multiple cases, the earliest generated case takes precedence. Other than that, execution order is free; in manual simulation mode, the user gets to choose, while automatic simulation mode chooses randomly.

The state of execution at any particular point is completely described by the state of all tokens: for each token its state is described by its location (a place, XOR, or transition), the executing role (if any), and the remaining processing time (if being processed by a timed transition). Each execution step either puts a token inside a transition or collector (consuming 0 or more tokens from places) or releases one from a transition or emitter (producing 0 or more tokens into places).

1.6.4.5 Case sensitive proper completion

When a cased token arrives at a collector, we say the collector *collects* the case. If no other token for the case remains, we say the case *completes*. A model in which every possible generated case is guaranteed to complete has *proper completion*.

In nets with proper completion, cases cannot continue forever, and cannot terminate without completing. Proper completion does not rule out the possibility of cases being collected multiple times.

1.7 Token color

In *colored Petri nets*, tokens can have values that can be tested and modified within the net. For instance, a transition can be specified to consume only integer tokens with a value smaller than 10, and to increment the value on tokens consumed.⁸

Color allows data to be associated with a process model.

The *token case* feature (section 1.6.4) can be regarded as a specific, limited application of token color. Yasper does not presently support any other use of color.

⁸ Process modeling tools such as CPN Tools (<http://wiki.daimi.au.dk/cpntools/>) and ExSpecT (<http://www.exspect.com/>) are based on colored Petri nets.

1.8 Data stores


A natural way to associate data with a process is to say that the data reside in a data store, and the process affects the data in a particular way. A typical example is a process interacting with a relational database by issuing SQL query statements.

Yasper allows such data stores to be represented explicitly.⁹ A store is like a place in that it is connected with tasks (transitions or XORs).¹⁰ For arcs connected to stores, the normal arc types do not apply, since tokens cannot be added to or removed from a store. Instead, it can be specified whether the task can

- create
- read
- update and/or
- delete

(the) value(s) in the store. Any combination can be specified, as long as at least one of the options is used.

Specifying this information (known collectively as a *CRUD matrix*) is a well-established technique; it allows basic consistency checking on the use of data.¹¹



In this example, *pay* modifies the transaction database, while *inspect overview* reads it, but does not modify it.

The present version does not use the specifications of stores and store arcs in any way; simulations ignore them. A future version of Yasper can support full details in this specification through a plug-in mechanism (see previous section).

⁹ The store concept was taken from ExSpecT.

¹⁰ In colored Petri nets, the contents of the store can be specified exactly as a token value.

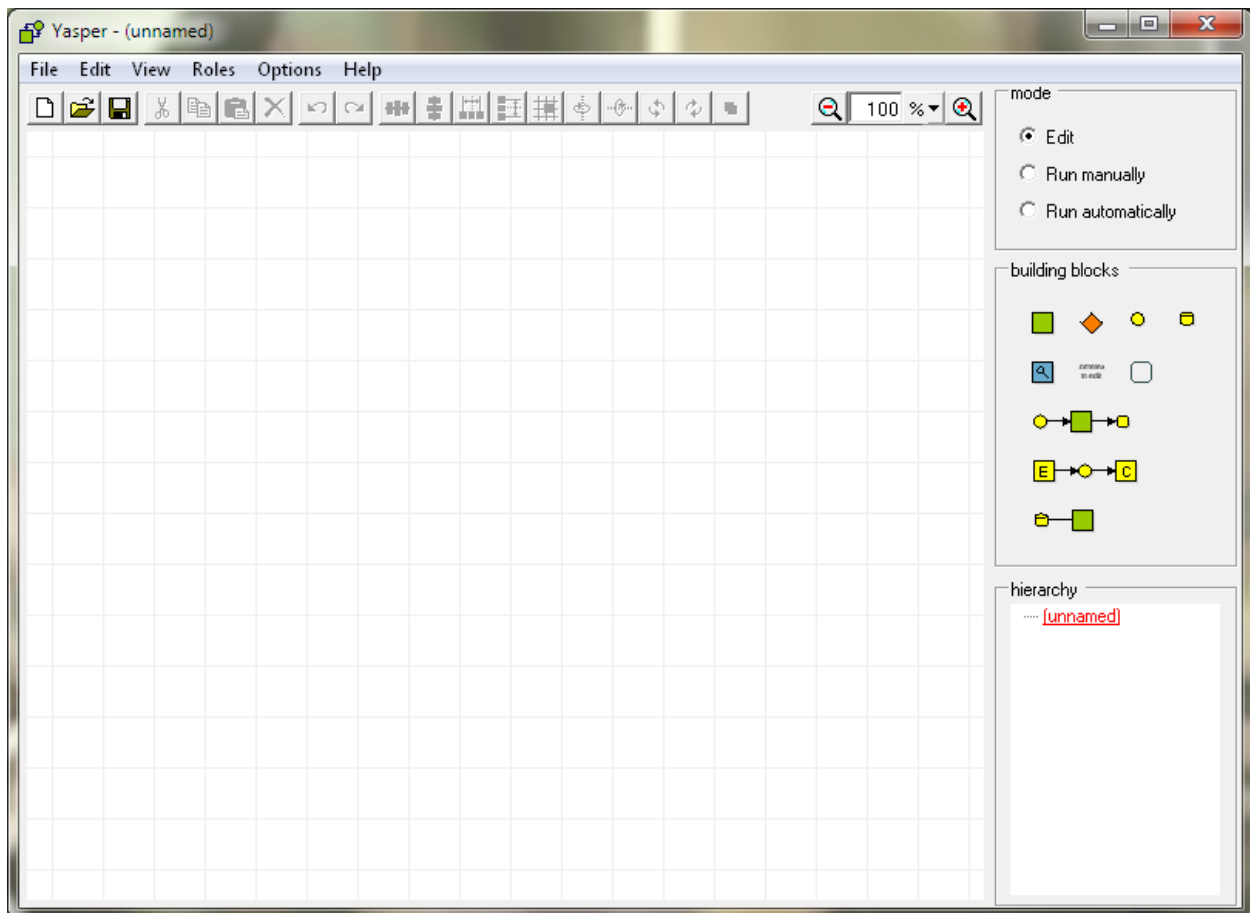
¹¹ Yasper does not provide any direct support for such checks.

2 The Yasper editor

This chapter provides a command-by-command introduction to Yasper's edit mode, which is used to create and modify process models.

2.1 Starting with an empty model

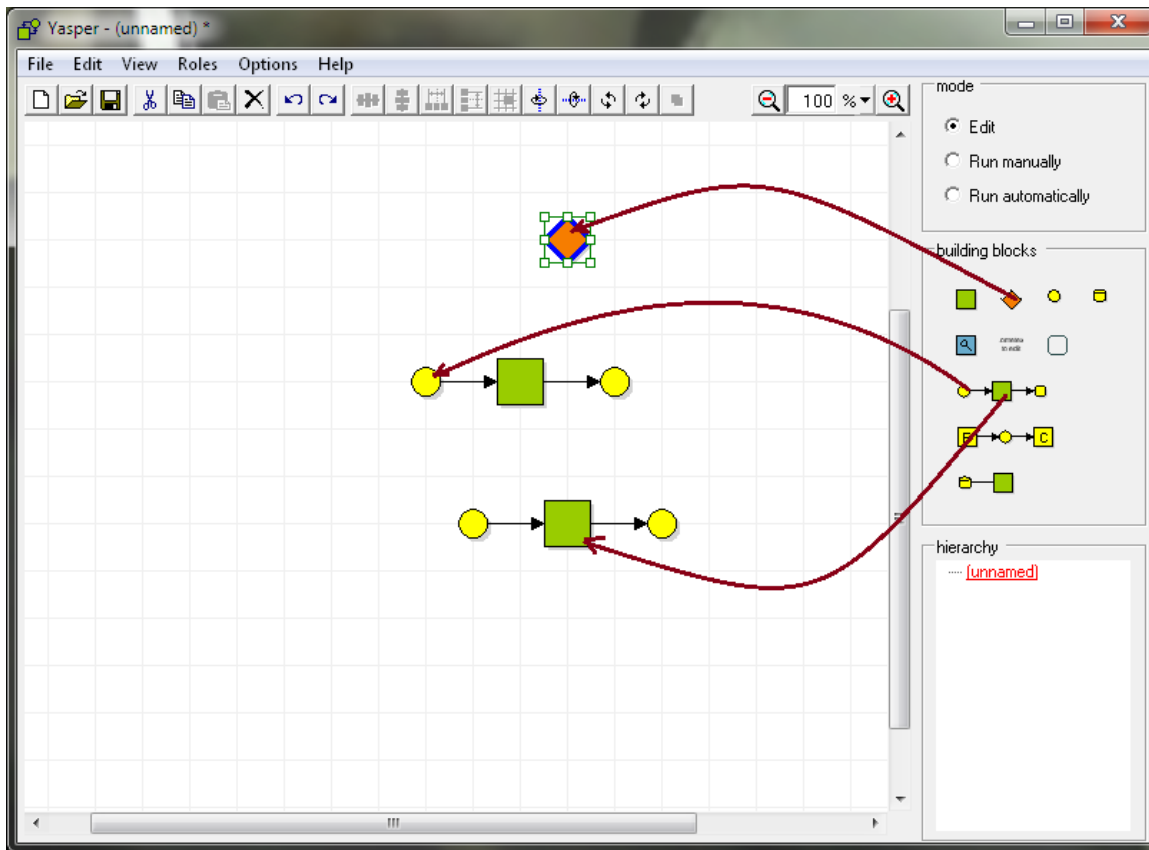
Start Yasper from the Windows Start menu or by double-clicking on the Yasper desktop short-cut. The main Yasper window comes up in edit mode.



A quick overview of the interface elements shown and their purpose in editing:

- the diagram canvas: to modify the model with direct manipulation
- the *File* menu: to load and save models in various file formats
- the *Edit* menu: to execute certain modifications (e.g. alignments) and copy/paste
- the *View* menu: to configure how models are displayed
- the *Roles* menu: to specify roles (they are not shown in the diagram)
- the *Options* menu: to configure how the editor behaves
- the *Help* menu: to offer and configure some basic assistance and link to this manual
- the toolbar icons: to directly execute many *File* and *Edit* menu operations
- the mode buttons: to switch between editing and simulation modes
- the building blocks: to add elements to the diagram
- the hierarchy view: to navigate and operate on the subnet hierarchy


Initially, the model is empty. To start editing, pick up a building block with the mouse and drop it onto the diagram canvas:



The scrollbars always allow all parts of the diagram to be centered on, regardless of its size.

2.2 Selecting and deselecting elements

All diagram manipulations work on a *selection* of elements; the selected elements are shown with blue borders. Additionally, a green box with handles is shown around the selection.





To select an element, left-click on it.

To add an existing element to the selection, shift-click on it.

To remove an element from the selection, shift-click on it.¹²

To select all elements in the diagram, choose *Select All* in the Edit menu, or press *Ctrl-A*.

To select a range of elements, click in the background, and drag the mouse across. A gray rectangle appears; the overlapping elements will be selected.



To add these elements to the existing selection instead, hold *Shift* while dragging.

To clear the selection, click in the diagram background.


Once a selection of elements has been made, it can be operated on, either by invoking a menu item from the Edit menu or clicking on an icon in the toolbar, or directly in the diagram by dragging or clicking on one of the selected elements. This element is the *primary selection*; it has a special role in some operations.¹³

¹²Most Windows applications use Ctrl-Click for this functionality; Yasper uses Ctrl-Click to connect nodes.

¹³Yasper's diagram editing functionality was shamelessly inspired by Microsoft Visio. The first Yasper prototype was implemented as a set of Visio macros.

2.3 Moving elements

To move the selected elements, click on one of them and drag the mouse in the desired direction.




To copy the selection and move the copy, hold down the *Ctrl* key while dragging.





Moves can also be executed with the arrow keys. This allows precise placement.



Alternatively, pick up the selection by a handle of the green selection box. This will not only move or copy it, but also change the relative position of its elements:



To constrain movement horizontally or vertically, pick a handle on the middle of a side instead of one on a corner:



If *Nodes are resizable* (in the *Options* menu, see section 2.18), the nodes in the selection will also change in size:





This is not recommended: Yasper does not offer any means to ensure that nodes of the same type will consistently be of the same size.

(A *node*, in this manual, is every element in the diagram that is not an *arc* – for instance, a *place* or a *transition*. Contrary to arcs, nodes have a position, height and width; as a consequence, many operations treat nodes and arcs differently.)

Arcs always stay connected.

Selecting multiple arcs, but no other elements, is impossible; it would not be very useful.

Trying to move a single arc will create a *support point* for it (see section 2.7):



After moving elements, automatic repositioning may occur:

- overlapping nodes may be *merged* (see next section), if *Merge nodes* is enabled in the *Options* menu;
- nodes are positioned on the grid, if *Snap to grid automatically* is enabled in the *Options* menu;
- nodes overlapping with unselected nodes cause the selection to be moved away , if *Fix overlaps automatically* is enabled in the *Options* menu.


The diagram does not have fixed boundaries: it will expand as needed, in all directions. However, Yasper becomes slow on diagrams with hundreds or thousands of elements; use *subnets* (see section 2.11) to keep diagrams manageable.

Dragging and dropping selections between different instances of Yasper is not supported. (Cut and paste is: see section 2.5.)

2.4 Merging elements

Moving a node onto another will fuse them together, under the following conditions:

- the resulting diagram makes sense
- the node is the element being pointed at (the primary selection)
- node merging is enabled in the *Options* menu



The node's properties will be those of the receiving node.¹⁴

Arcs may be merged as a side effect of nodes being merged.

¹⁴ Except whether a transition is *emitter* or *collector*, a property determined by how it is connected to places.

2.5 Deleting, copying and pasting elements

To delete the selected elements, either

- select *Delete* from the *Edit* menu
- click the *Delete* toolbar icon
- right-click it and select *delete* from the pop-up menu, or
- press the *Delete* key.

Deleting a node also deletes all arcs attached to it – arcs must always be connected on both ends.

Alternatively, *cut* the elements into the Windows clipboard: either

- select *Cut* from the *Edit* menu
- click the *Cut* toolbar icon
- right-click it and select *cut* from the pop-up menu, or
- press *Ctrl-X*.

Likewise, four options exist to *copy* elements to the Windows clipboard.

Cut is effectively *copy* followed by *delete*.

If the selection contains any arcs with unselected endpoints, these endpoints are placed on the clipboard as well. This ensures that no loose arcs will result from pasting.



To *paste* the elements in the clipboard into a Yasper diagram, either:

- select *Paste* from the *Edit* menu
- click the *Paste* toolbar icon
- right-click on the diagram background and select *paste here* from the pop-up menu
- press *Ctrl-V*.

In all cases except the third, the contents will be pasted to the center of the diagram.

2.6 Connecting nodes

To create connections from the selected node(s) to another node, *Ctrl*-click on that node, which must not be part of the selection (if it is, the node is copied instead, cf. section 2.3).




Only valid connections will be made.

Intermediate nodes are inserted as required.

.

2.7 Extending the node frontier

To add “the next nodes” relative to the currently selected ones, drag the selection while pressing the *Alt* key:

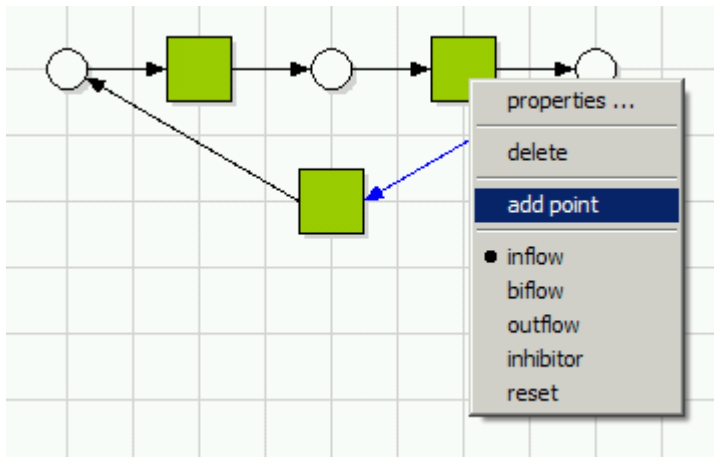


This operation can be very useful¹⁵ because elements of process models are usually created roughly in the order in which they are typically executed.


¹⁵ In this case, the inspiration was another Petri net tool: WoPeD.

2.8 Arc support points


To add an intermediate point¹⁶ to an arc, either drag on it (see section 2.3) or right-click on it and select *add point*.



To move a support point, left-click on it and drag.



Arbitrarily many support points can be added to an arc.

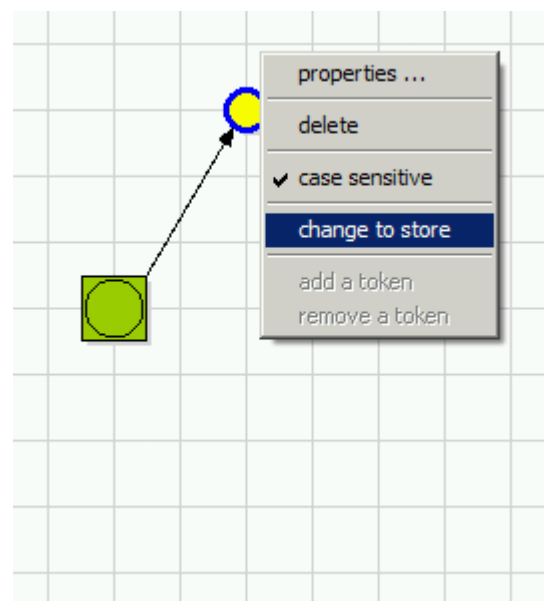
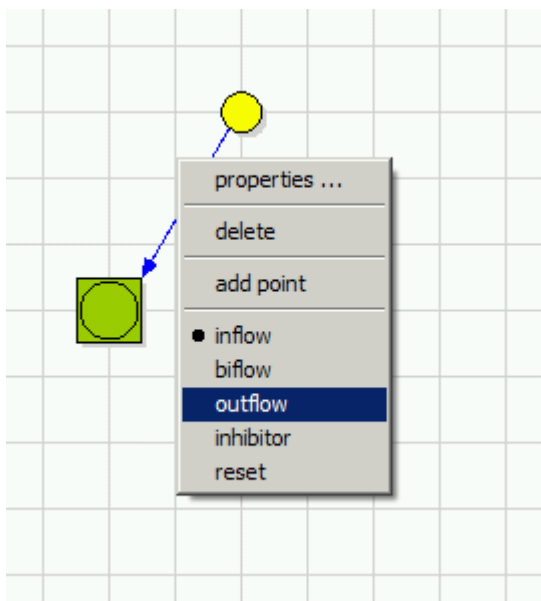
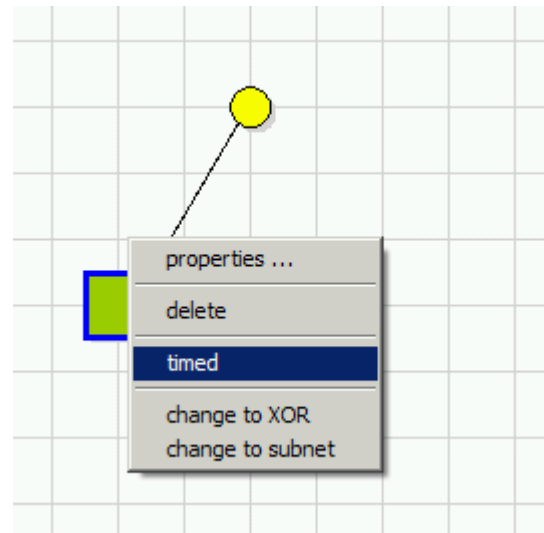
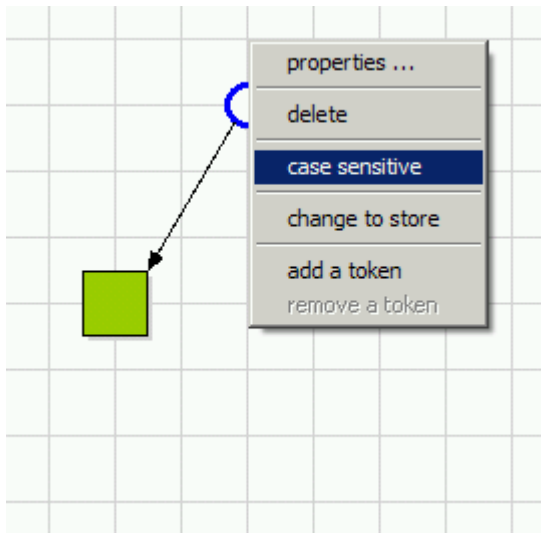


To delete a support point, right-click on it and select *delete point*.

¹⁶ No standard term appears to exist for the intermediate points of polylines; we use *support point*.

2.9 Changing the type of an element

To change the type of an element, right-click on it and select the new type.



Changing the type of a node automatically adjusts the types of the attached arcs, when necessary.

The available node and arc types and their meanings are discussed in chapter 1.

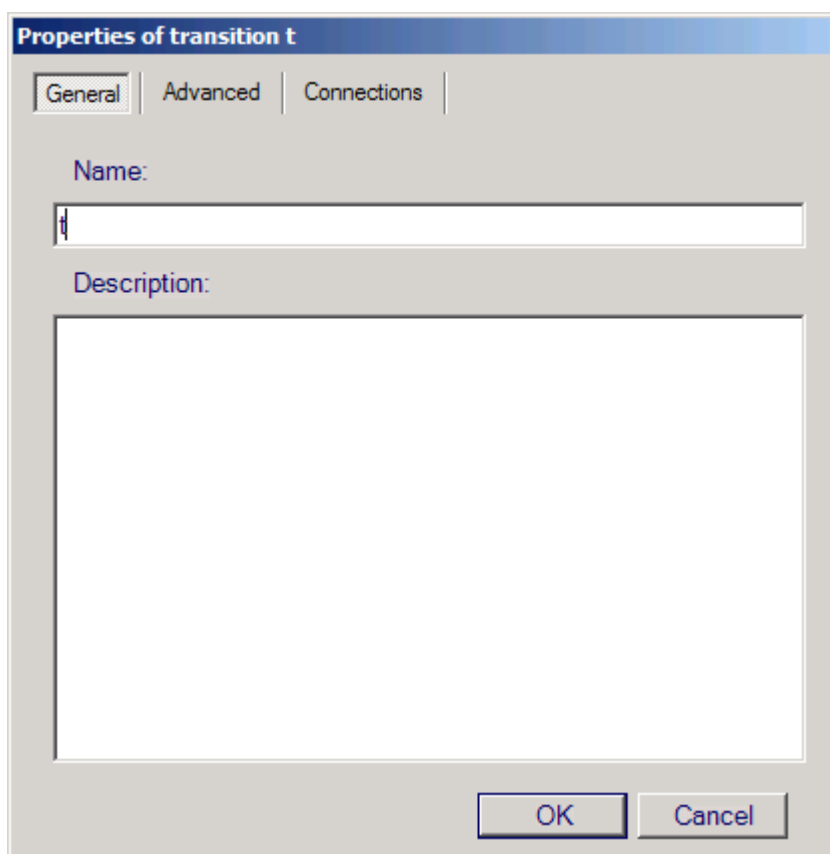
2.10 Specifying additional properties

Model elements have additional properties that can be set and inspected in the element's property sheet.

To bring up the property sheet for an element, right-click on the element and select *properties ...*, or double-click on it (except on subnets and pins).

To bring up the property sheet for the current net, right-click in the background and select *properties ...* or double-click in the background.

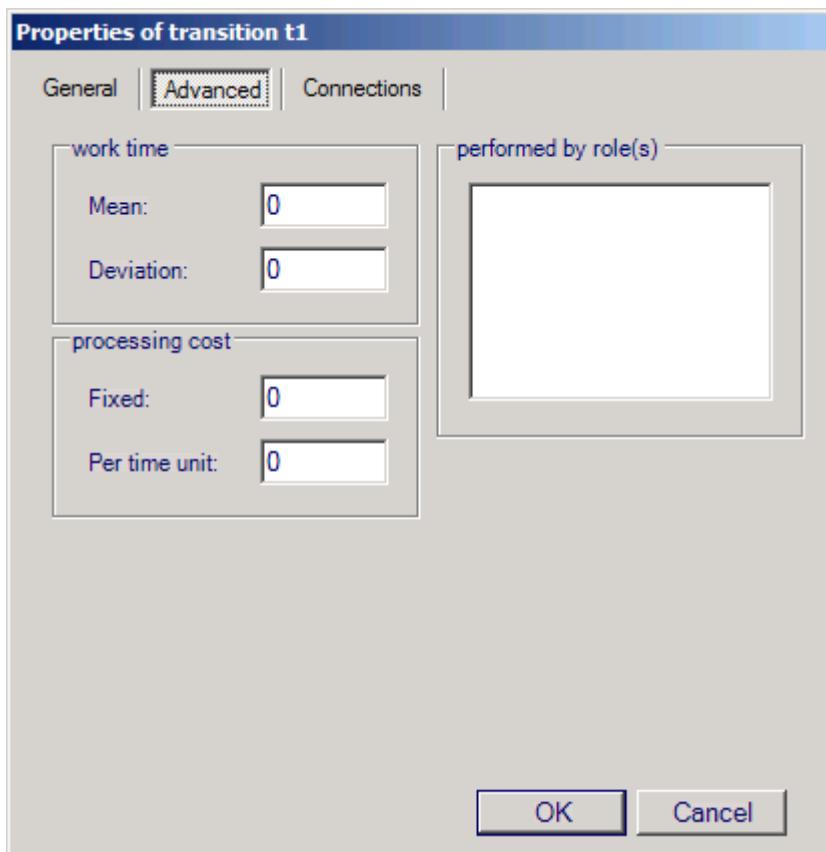
Every element has a name and a description. They can be entered in the property sheet. The name will appear in the diagram if *Show names* is enabled in the *View* menu. Names are optional and do not have to be unique.



Some elements have additional properties. To inspect or modify them, click on *Advanced* and fill out the relevant form field(s). You will be prevented from entering invalid information.

The various advanced properties are the following.

2.10.1 Transition properties



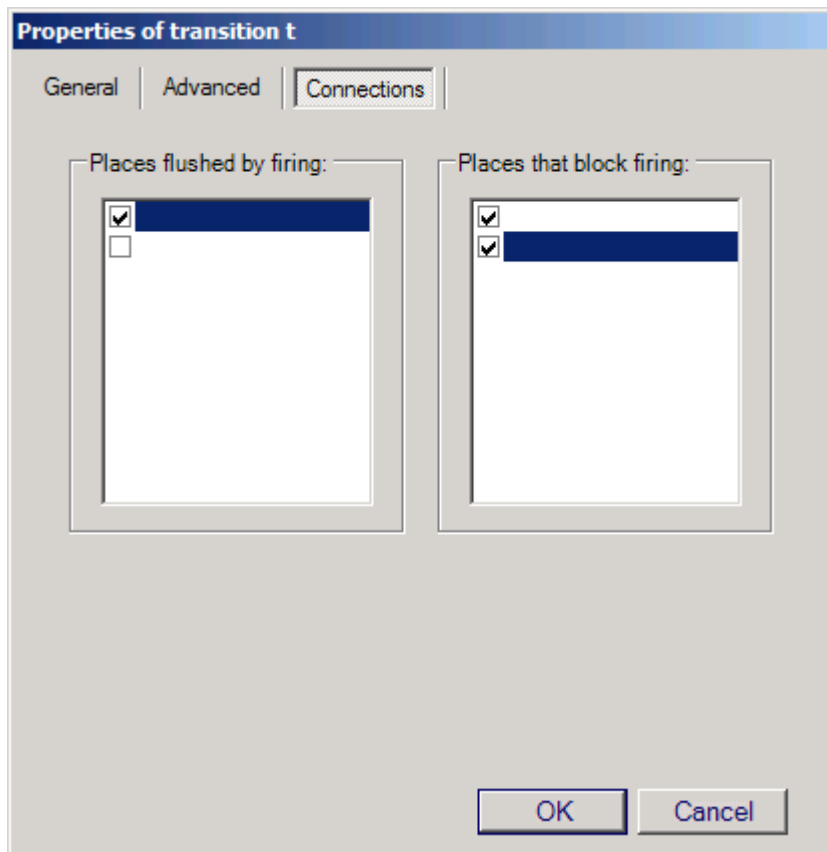
The image shows a dialog box titled "Properties of transition t1". It has three tabs: "General", "Advanced", and "Connections". The "Advanced" tab is currently selected. Inside the "Advanced" tab, there are two main sections. The first section is labeled "work time" and contains two input fields: "Mean:" with a value of 0 and "Deviation:" with a value of 0. The second section is labeled "processing cost" and contains two input fields: "Fixed:" with a value of 0 and "Per time unit:" with a value of 0. To the right of these sections is a large empty rectangular box labeled "performed by role(s)". At the bottom right of the dialog box are two buttons: "OK" and "Cancel".

For a transition, the parameters for processing time and cost (see sections 1.6.2 and 1.6.3) can be entered here. They must be decimal numbers.¹⁷

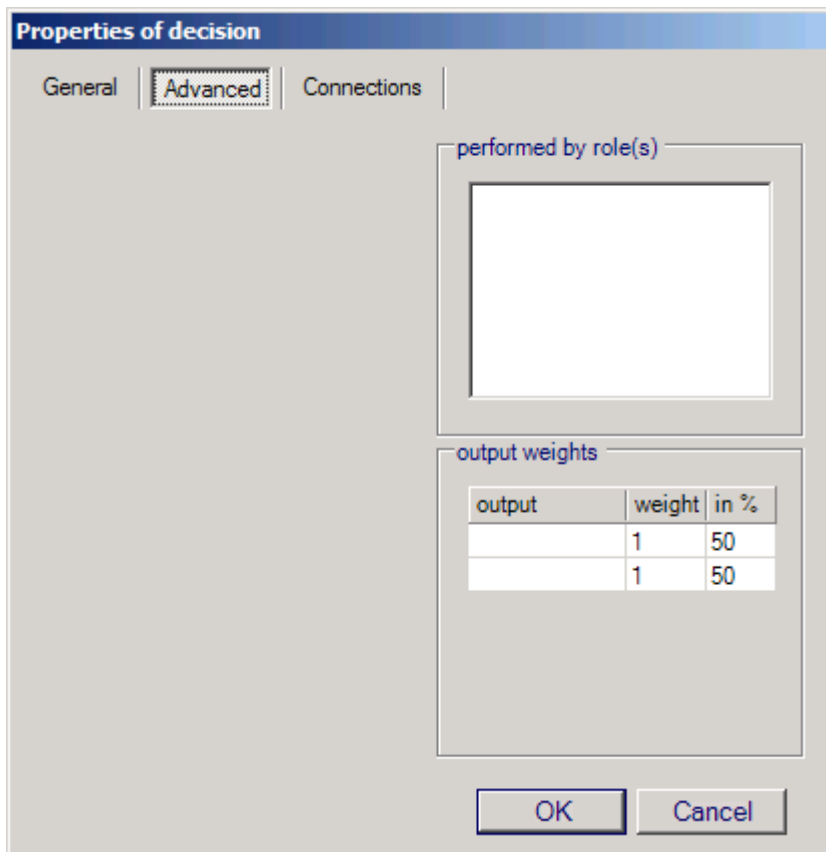
For role assignment see section 2.12.2.

The *Connections* tab can be used to add and remove reset and inhibitor arcs. Whether the arcs in question appear in the diagram is configurable in the *View* menu.

¹⁷ Caveat: even though Yasper's interface language is English, the decimal separator to use here is the one set in your Windows settings. This may change in the next version of Yasper.



2.10.2XOR properties



XORs allow relative output weights to be set (see section 1.6.1).
Like processing time and cost, output weights can have decimal separators.

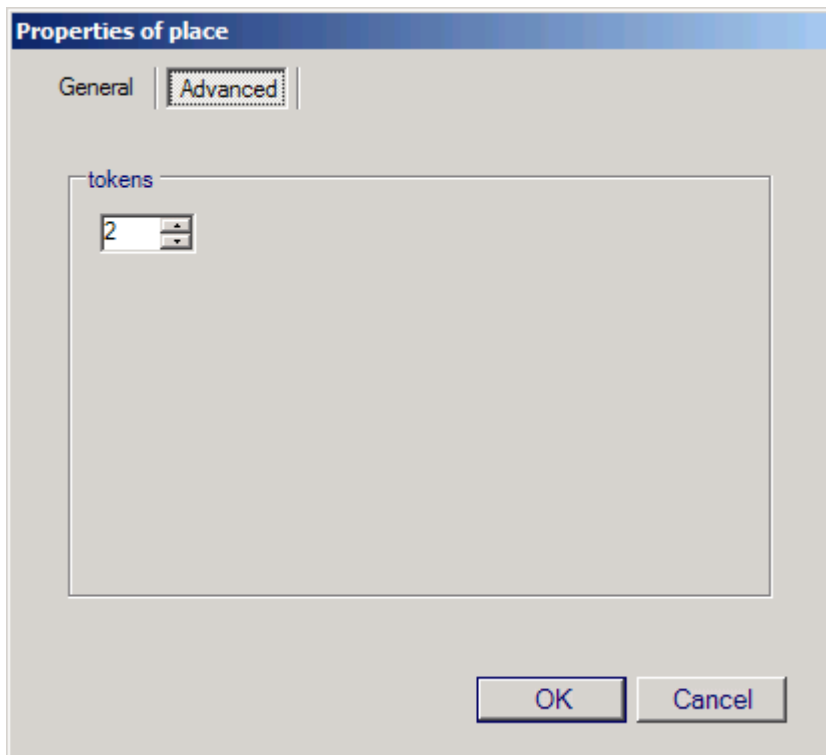
Type or edit a number in the *weight* column to adjust it; the percentages are recomputed automatically.

Type or edit in the output column to adjust or set the name of the place in question.

2.10.3 Place properties

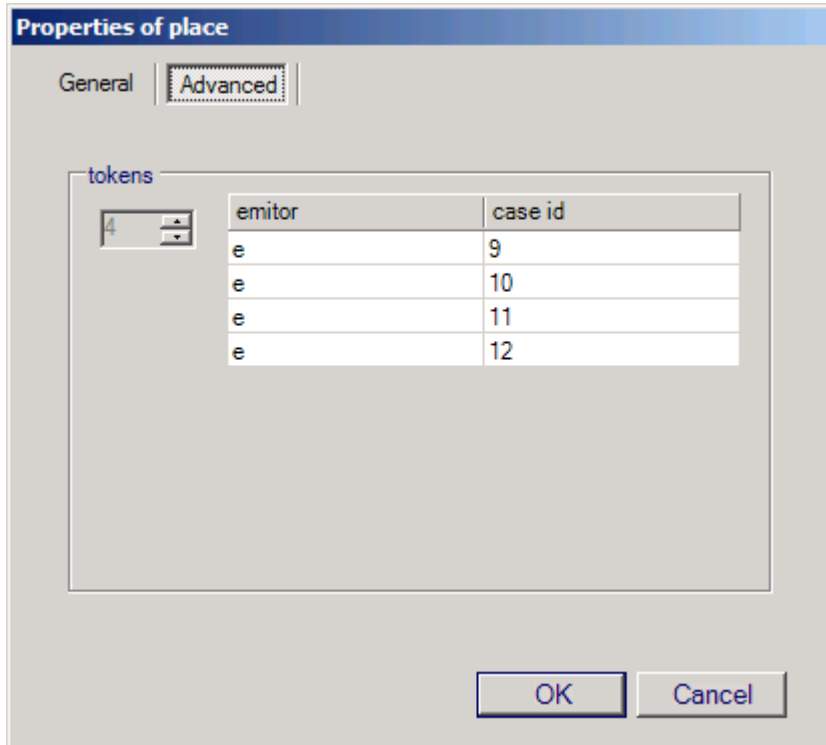
For places, the *Advanced* tab displays their tokens.

To change the number of tokens in the place, type it directly, or click on the tiny triangles next to it. Tokens do not carry values, other than the case for cased tokens – see section 1.7.



The dialog box titled "Properties of place" has two tabs: "General" and "Advanced". The "General" tab is selected. Inside the dialog, there is a label "tokens" above a text input field containing the number "2". At the bottom right, there are "OK" and "Cancel" buttons.

For case sensitive places, tokens are displayed, but cannot be edited.
The case of tokens is shown: emitter name and case sequence number.

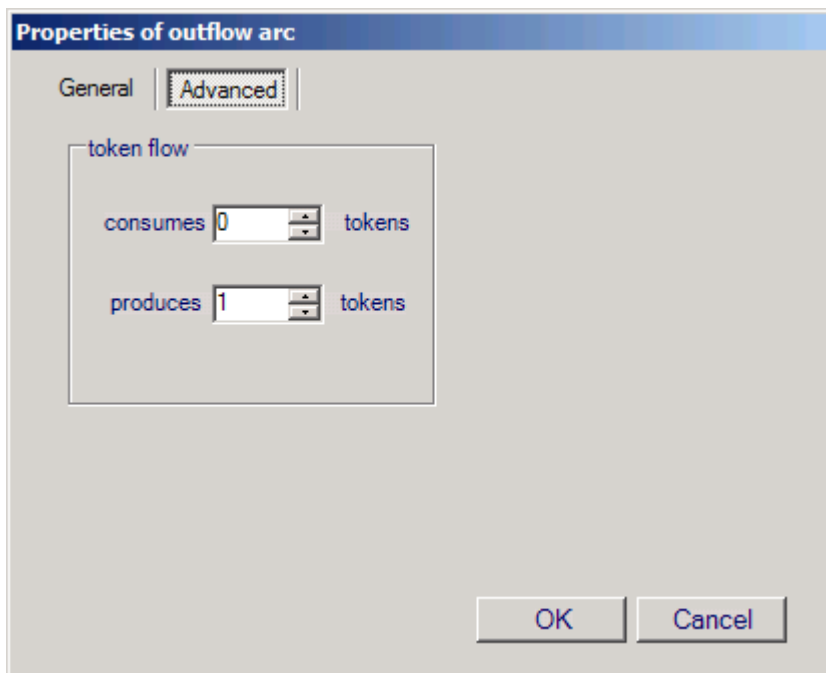


The dialog box titled "Properties of place" has two tabs: "General" and "Advanced". The "Advanced" tab is selected. Inside the dialog, there is a label "tokens" above a text input field containing the number "4". To the right of the input field is a table with two columns: "emitor" and "case id". The table contains four rows of data. At the bottom right, there are "OK" and "Cancel" buttons.

emitor	case id
e	9
e	10
e	11
e	12

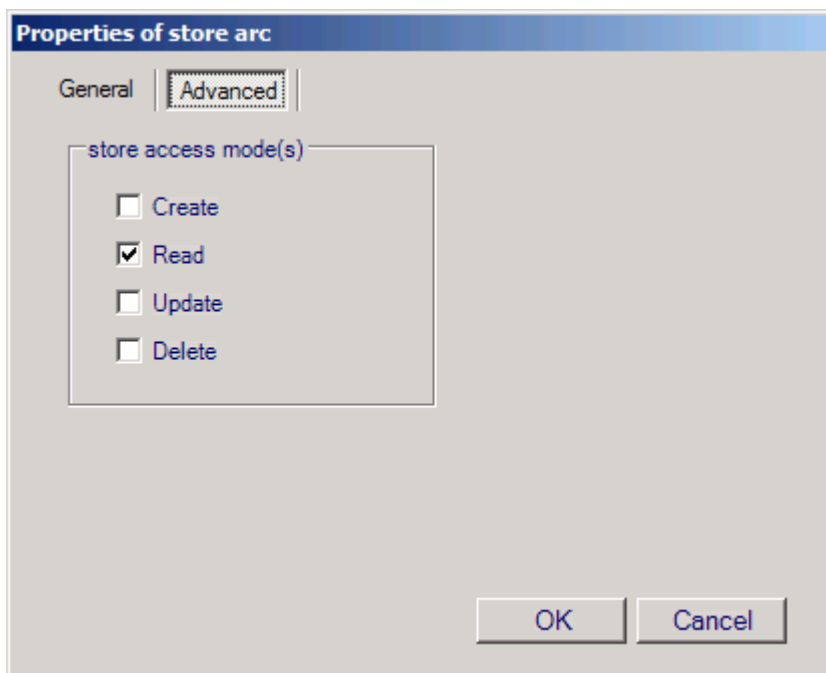
2.10.4 Arc properties

To set the number of consumed and/or produced tokens on a flow arc, open the *Advanced* tab of its property sheet and adjust the numbers shown.



The numbers cannot both be 0. With the numbers set to *1,0* or *0,1*, the arc is an inflow or outflow, respectively; otherwise, it is a *combined arc* (see section 1.5.1).

On arcs attached to stores, the kind(s) of store access performed can be set:



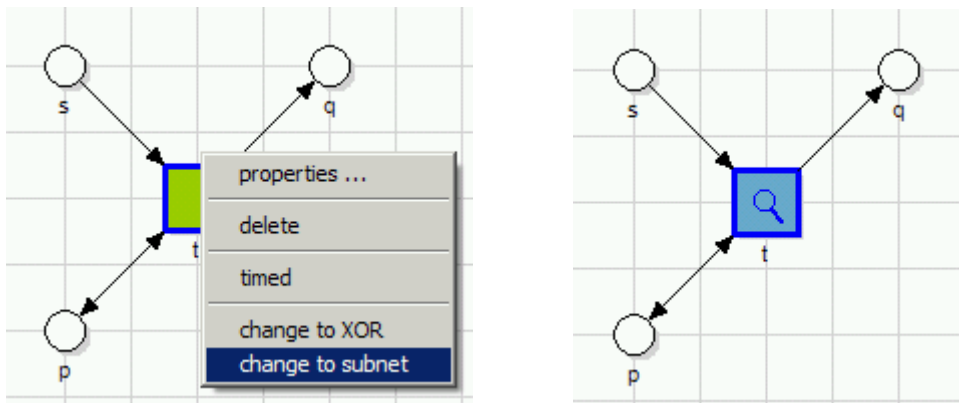
At least one must be selected.

2.11 Working with subnets

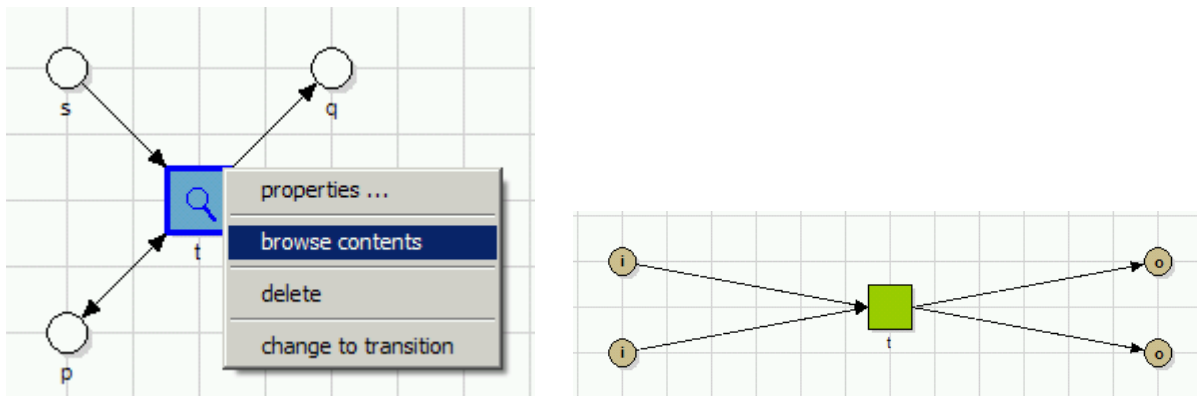
Subnets distribute models across multiple screens; full details in section 1.4.

2.11.1 Creating, navigating, and removing subnets

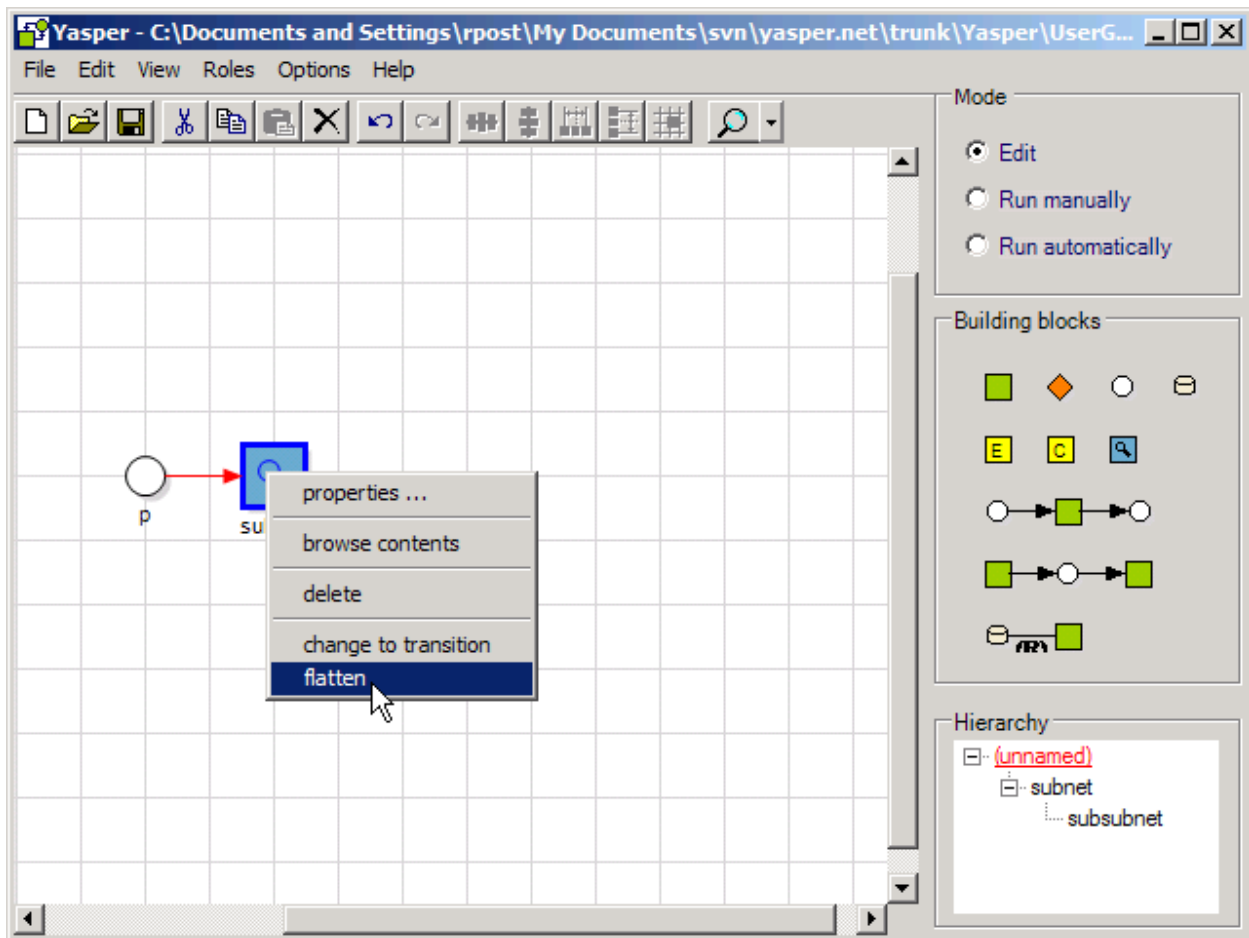
To convert a transition to a subnet, right-click on it and select *change to subnet*.



To move into the subnet, double-click on it, or right-click and select *browse contents*.



To move up to the surrounding (sub)net, double-click on a pin.



To remove a subnet, right-click on it and select either of three options:

- *delete* removes the subnet with its connections and all of its contents;
- *change to transition* replaces it with a transition, removing all of its contents;
- *flatten* replaces it with its contents, pushing the rest of the diagram out of the way.

2.11.2 Connecting across subnet boundaries

Transitions inside a subnet can be connected to places outside. To show such a connection in the subnet content, a *pin* exists within the subnet for every arc attached to the subnet in the surrounding net.


Pins can be manipulated like places, with two restrictions.¹⁸ First, pins cannot be copied or merged. Second, at most one arc can be attached to a pin, directly or indirectly; attempts to attach another will fail.

To add a pin to a subnet, connect the subnet to a pin or place.

¹⁸ These restrictions will be removed in a future version of Yasper.



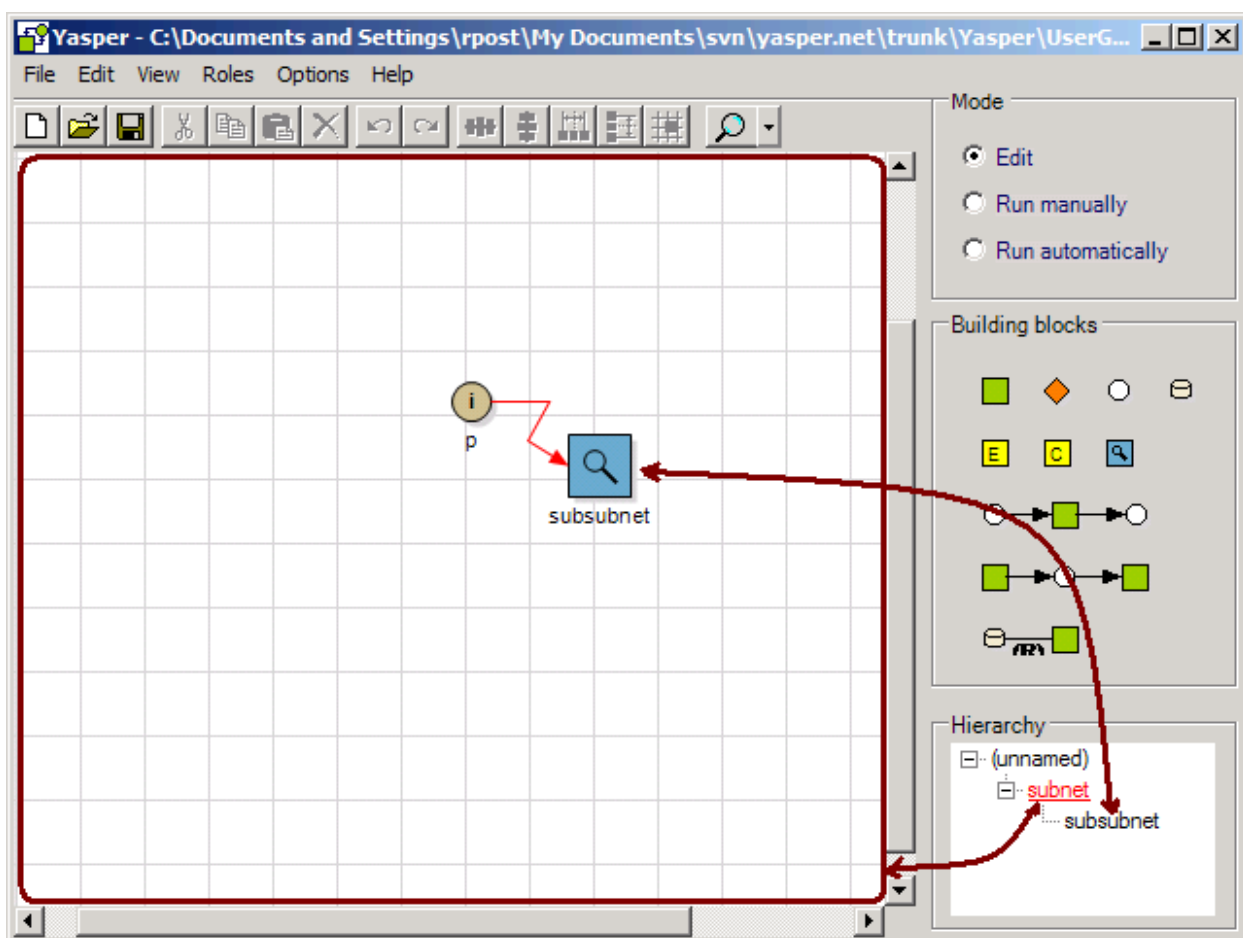
An arc on a subnet is shown in red whenever its pin inside is not connected to a transition, directly or indirectly. For example, in



creating an arc between t and the pin will turn the two red arcs black.

2.11.3 Working with the tree view


At the bottom right, Yasper displays a tree view of the subnet hierarchy. The last example, with a subnet inside a subnet, is displayed as follows.



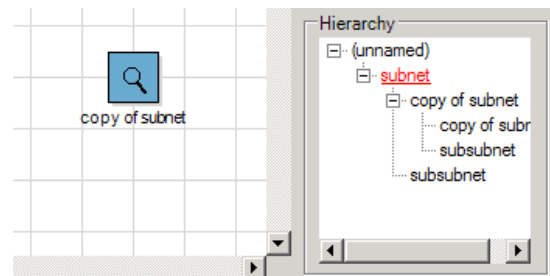

The net presently displayed is marked in the tree view: its name is in red and underlined.

To navigate to a (sub)net of choice, click on its item in the tree view.

To rename a net, click on its tree view item, wait, then click on it again.¹⁹ A rectangle will appear around it. Type or edit the name.



To copy a (sub)net into the current net, drag its tree view item onto the diagram canvas. Since nets may not directly or indirectly contain themselves, recursion is cut short.




¹⁹ You may have to click three times.

2.12 Working with roles

Roles provide a way to model resources required to execute a task (see section 1.3).

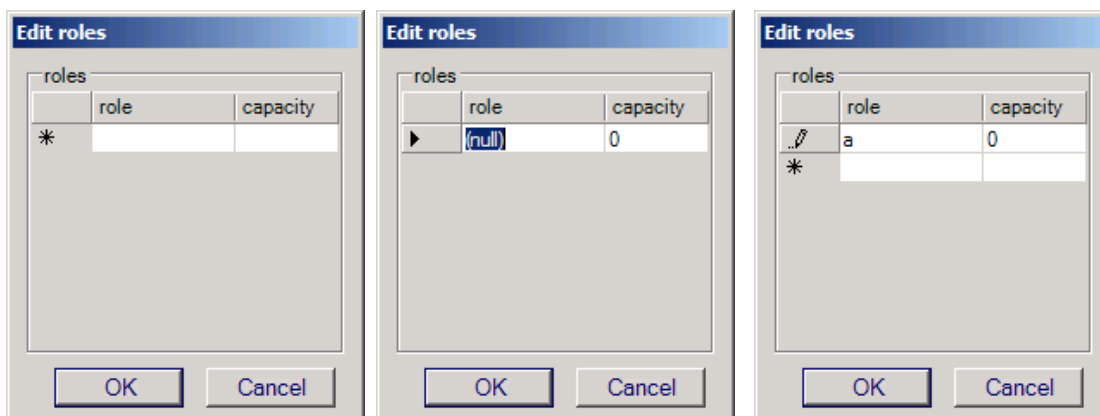
2.12.1 Defining roles

On the *Roles* menu, select *Define roles ...*



The *Edit roles* dialog is shown.

To add a new role, click in the *role* column and type a name for it.



To delete a role, click on the leftmost, gray part of its row and press the *Delete* key.

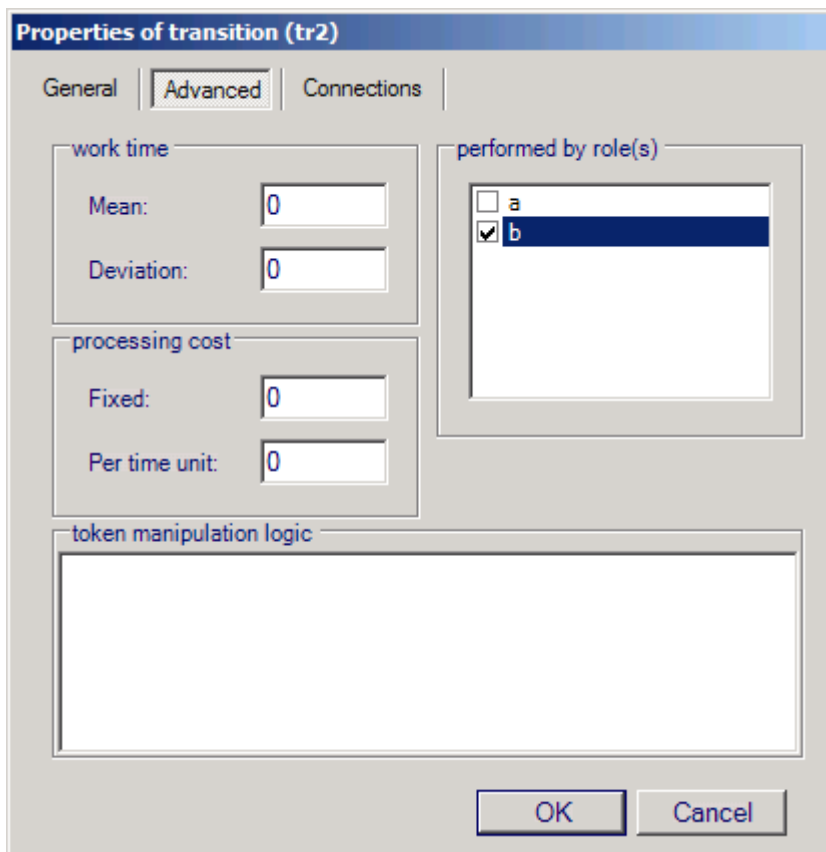
To set a role's capacity, type the number in the capacity column. A capacity must be a positive integer; *0* is invalid and will be automatically corrected to *1*.

Unlike the names of other elements, role names are mandatory and unique. Yasper automatically adjust the names entered to ensure this.

2.12.2 Assigning roles to transitions

Once one or more roles have been defined, they can be assigned to transitions and/or XORs.

To assign/deassign roles for an individual task (i.e., transition or XOR), double-click on it in the diagram, select its *Advanced* properties (see also section 2.10.1), and check the role(s) to be assigned.

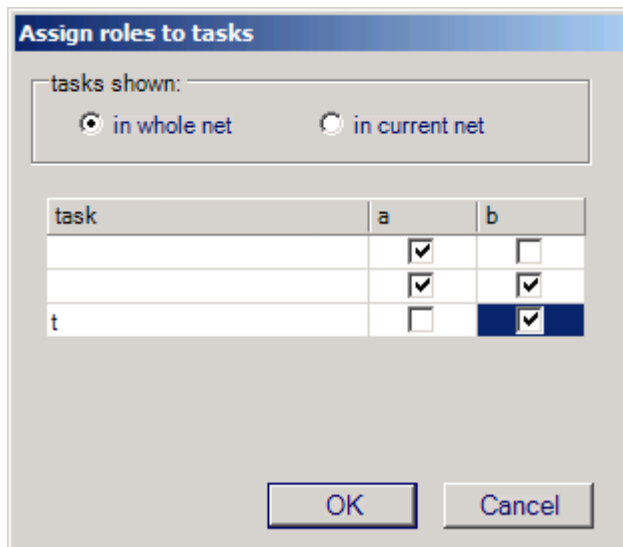


The dialog box titled "Properties of transition (tr2)" has three tabs: "General", "Advanced", and "Connections". The "Advanced" tab is selected. It contains three sections: "work time" with "Mean:" and "Deviation:" fields both set to 0; "processing cost" with "Fixed:" and "Per time unit:" fields both set to 0; and "token manipulation logic" with a large empty text area. On the right, the "performed by role(s)" section shows a list with roles "a" and "b", where "b" is checked. At the bottom are "OK" and "Cancel" buttons.

For an overview of all role assignments, select *Assign to tasks ...* in the *Roles* menu.



The role assignment dialog is shown. It has a column for each role and a row for each task (i.e. transition or XOR).



To assign a role to a task, check the relevant checkbox.

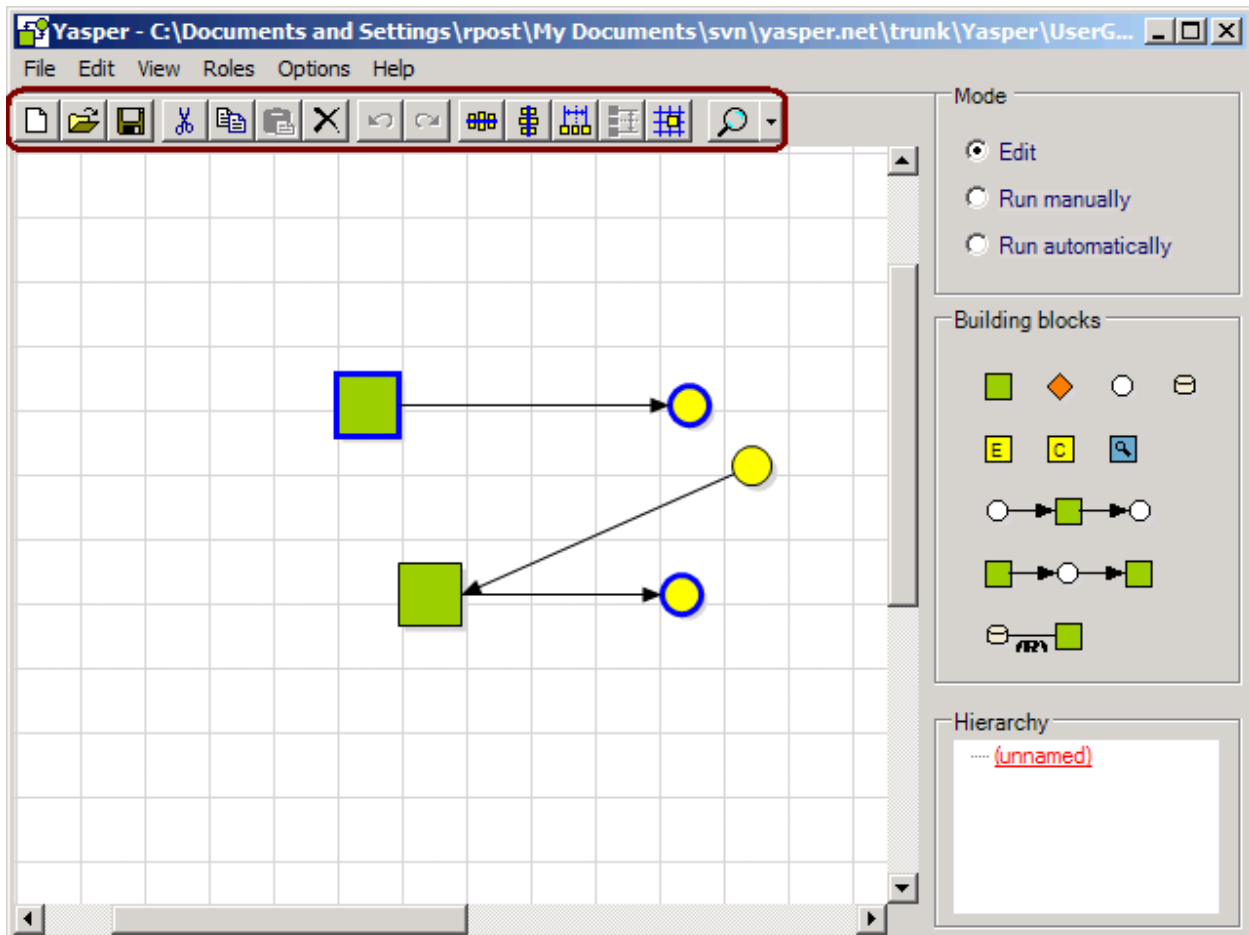
To deassign it, uncheck it.

To limit the rows to tasks in the (sub)net being displayed in the diagram, click on the *in current net* button.

If tasks have not been named, it may be practical to set *Show identifiers* in the *View* menu prior to bringing up this dialog.

2.13 Toolbar operations


The toolbar below the main menu items provides one-click access to frequent operations. The toolbar items are just shortcuts: each item corresponds to a menu item in the *File* or *Edit* menu, explained in sections 2.15 and 2.14, respectively.



An item is grayed out if the operation cannot be performed at the present time.

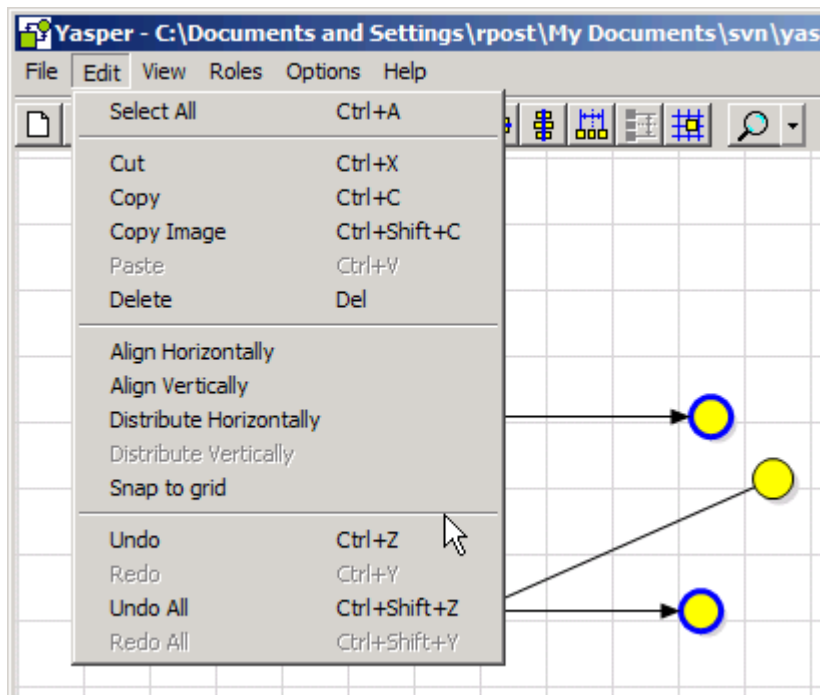
To execute an operation, click the item in question.

To get a textual description, put the mouse over an item and leave it still.



2.14 *Edit* menu operations

The *Edit* menu provides the standard operations and a few more.



2.14.1 Undoing changes

Yasper remembers all changes made in the editor since the last time the model was saved to file.

To undo a change, select *Undo* in the *Edit* menu, click the *Undo* toolbar icon, or press *Ctrl-Z*.
To redo a change, select *Redo* in the *Edit* menu, click the *Redo* toolbar icon, or press *Ctrl-Y*.

To *Undo* as many times as possible, the shortcut *Undo All* is available.
To *Redo* as many times as possible, *Redo All*.

2.14.2 Copying and pasting

To copy the elements selected in the diagram, select *Copy* in the *Edit* menu, click the *Copy* icon on the toolbar, or press *Ctrl-C*.

The copy is not always exact, since it is adjusted to form a valid model:


- if an arc is selected, the nodes it connects are in the copy, even when they were not selected;
- if a pin is selected, it is added to the copy as a place.

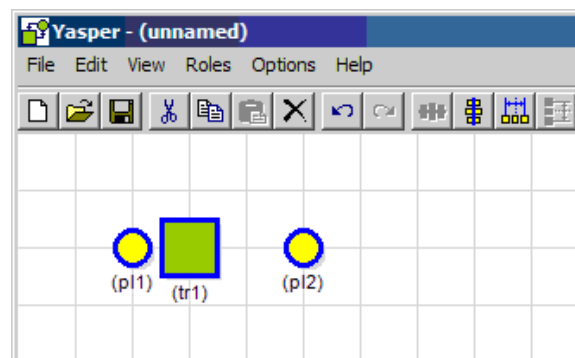
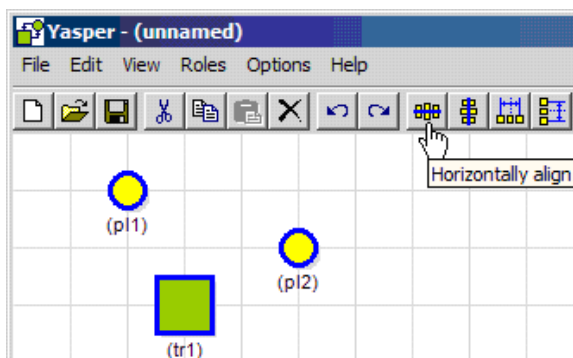
To paste a copy into the model, select *Paste* in the Edit menu, click the *Paste* icon on the toolbar, or press *Ctrl-V*. Pasting works in other (sub)nets and even in other Yasper instances, but not in other applications.


After pasting, the new elements are selected so they can be moved into position.

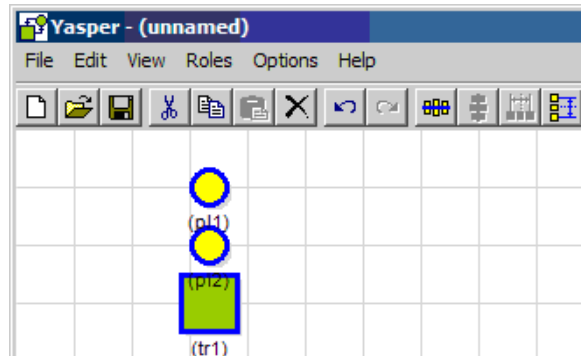
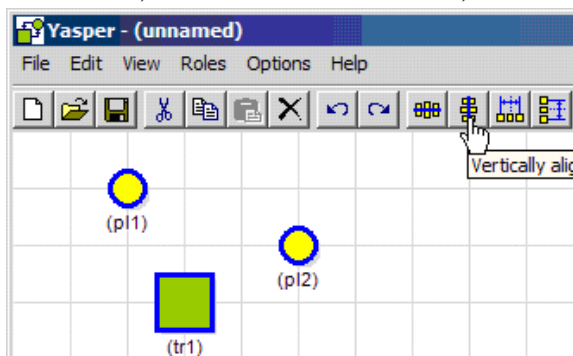
To copy an image of the current net that can be pasted into other applications,²⁰ use *Copy Image* on the *Edit* menu. The image will contain all elements of the current (sub)net, with the selected elements marked; gridlines are not shown.

2.14.3 Aligning elements

To shift the selected elements so they are on the same horizontal line, select *Align Horizontally* on the *Edit* menu, or click the toolbar icon, .



To shift the selected elements so they are on the same vertical line, select *Align Vertically* on the *Edit* menu, or click the toolbar icon, .




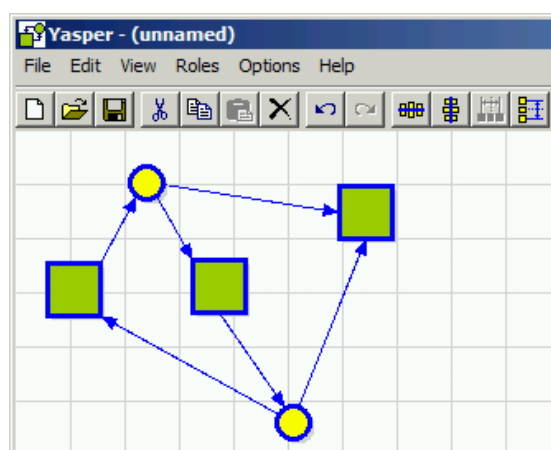
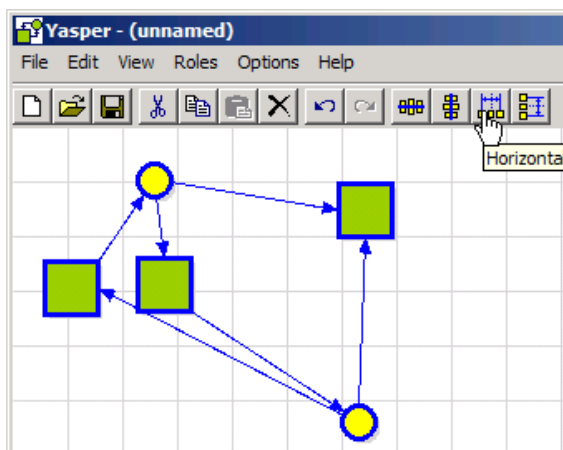
Arc support points (see section 2.7) are not affected by the operation.

Another way to keep node alignment tidy is to snap them to the grid - see section 2.14.5.

²⁰ The Microsoft Office clipboard sometimes refuses to accept an image; when this happens, clear the clipboard before issuing *Copy Image*.

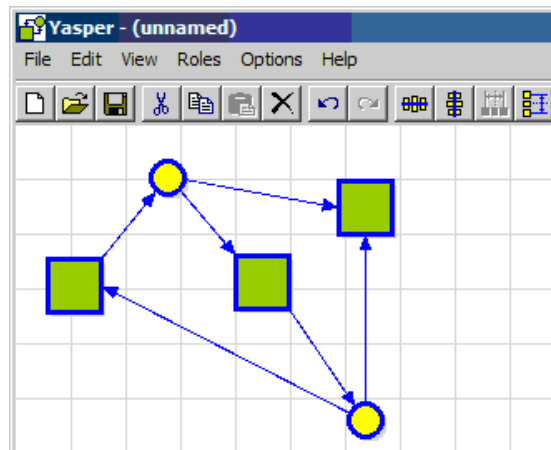
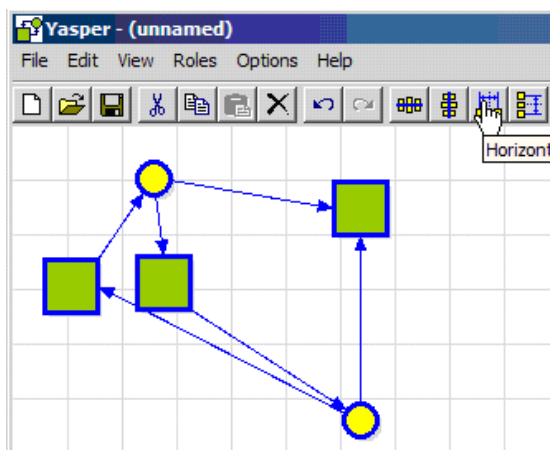
2.14.4 Making elements evenly spaced

When elements are in a horizontal sequence, it often looks tidier to distribute them to be evenly spaced. To perform this operation on the selected elements, select *Distribute Horizontally* on the *Edit* menu, or click on the toolbar icon, .




This works well to bring out the logical ordering between elements even when they are not actually horizontally aligned.

Nodes on the same height will remain on the same height:



This operation completely ignores arcs: arcs do not affect how nodes are displaced, and arc support points stay where they are.

Horizontal distribution is orthogonal to horizontal alignment: distribution shifts nodes horizontally, while alignment shifts them vertically. They are often combined.

To apply the same operation vertically, select *Distribute Vertically* on the *Edit* menu, or click on the toolbar icon, .

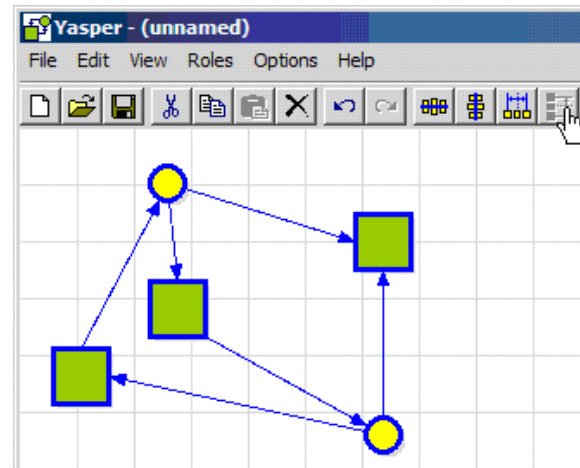
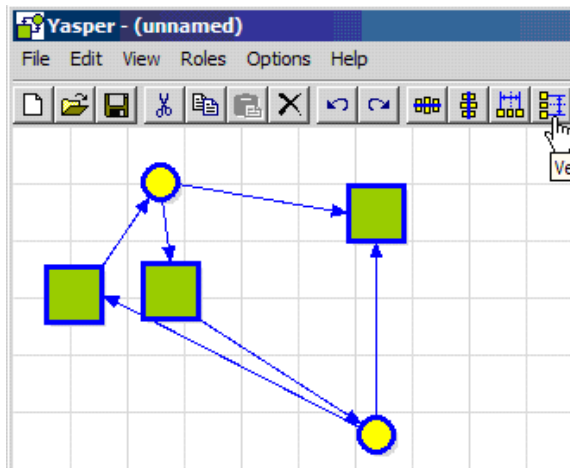



Diagram layouts can be tidied up efficiently and quickly by well-chosen combinations of horizontal and vertical alignment and distribution operations.

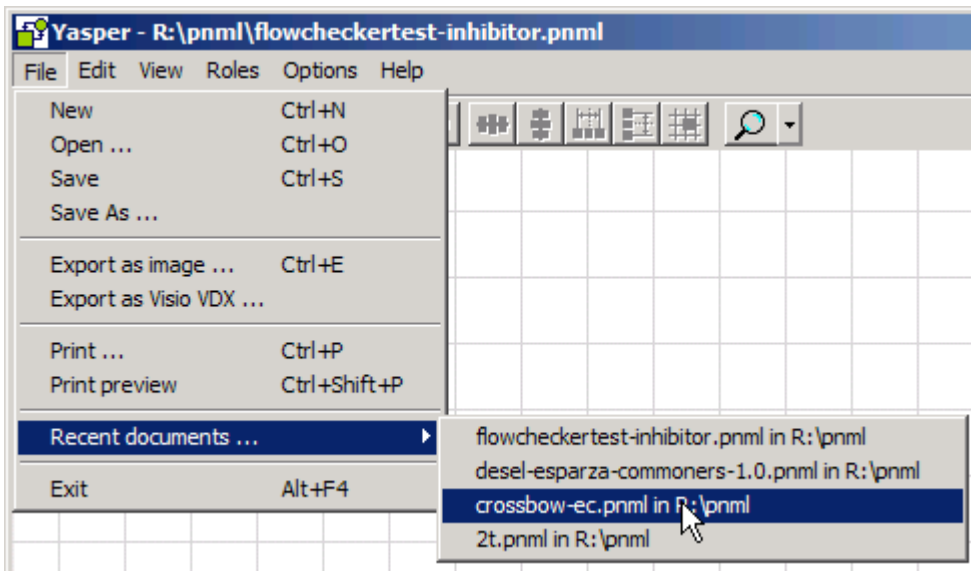
2.14.5 Snapping to the grid

Use the *snap to grid* menu item or  toolbar icon to move the selected nodes and arc support points to the nearest location on the grid. Working with the grid in this way is an alternative to the alignment and distribute operations described in the previous paragraphs.

The editor can be instructed to perform this operation automatically whenever a node or arc support point is created or moved, by setting the *Nodes snap to grid* option and/or the *Arc points snap to grid* option in the *Options* menu (section 2.18).

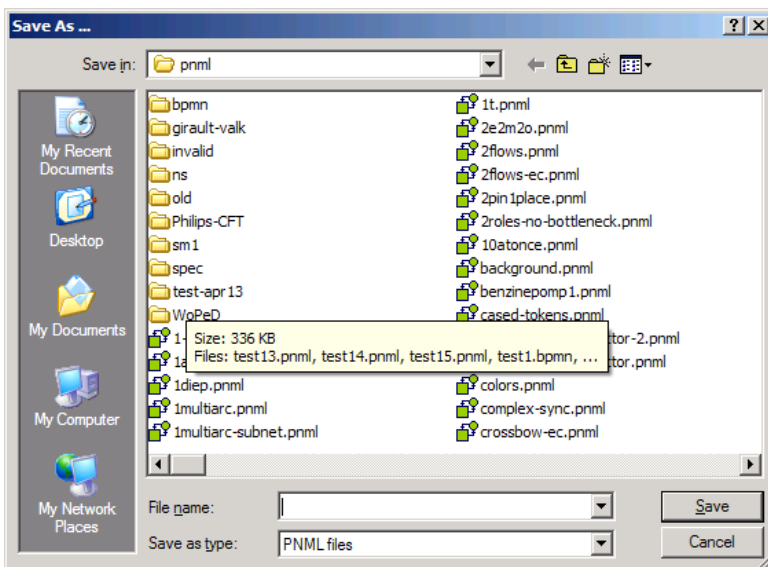
2.15 File menu operations

The Yasper *File* menu offers operations to load and save models, to export them to various file formats, to print images of models, and to exit Yasper.



2.15.1 Loading and saving Yasper models

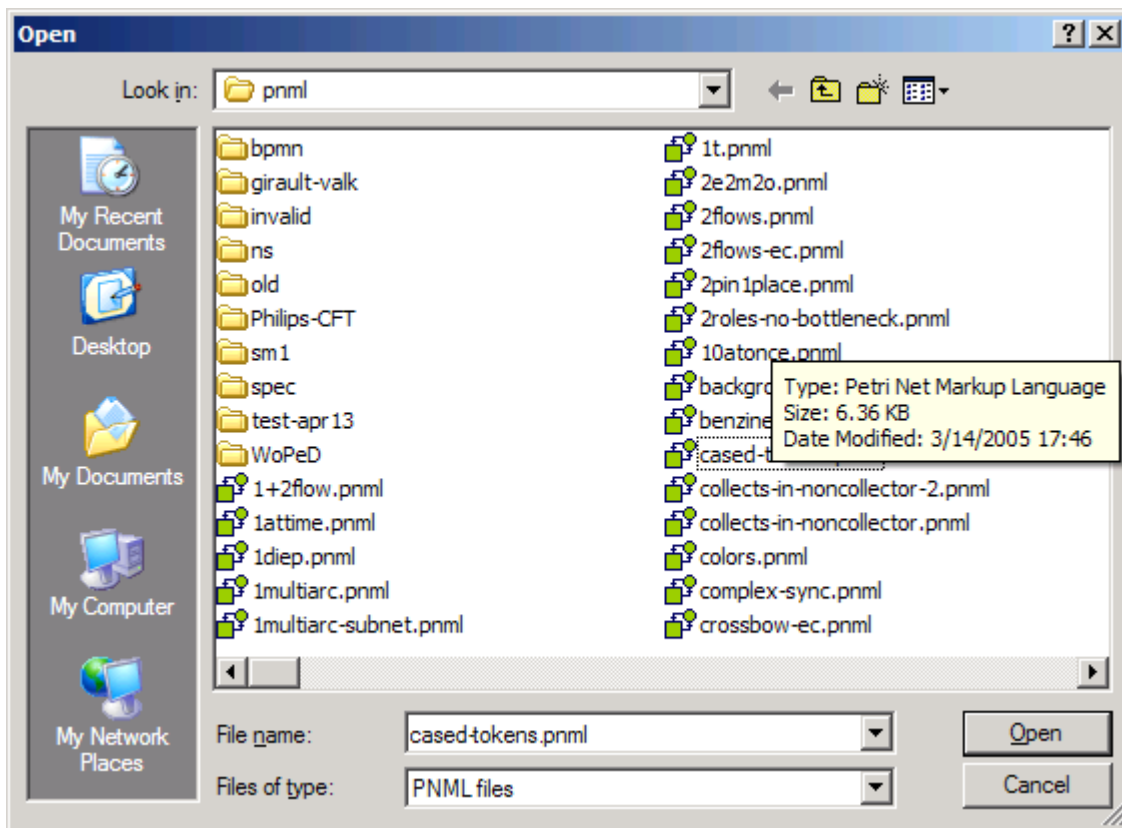
To save a model to file and specify the filename, select *Save As ...* on the *File* menu. A dialog will ask for the filename.



To save a model and use the filename with which it was opened or saved before, select *Save* on the *File* menu, click the *Save net* toolbar icon, or press *Ctrl-S*. If Yasper doesn't know the filename to use yet, it shows the *Save As* dialog.

To start working on a new model, select *New* on the *File* menu or press *Ctrl-N*.
If the existing model has unsaved changes, Yasper will offer the opportunity to save it first.

To open an existing model, select *Open* on the *File* menu, click the *Open net* toolbar icon, or press *Ctrl-O*. If the existing model has unsaved changes, Yasper will offer the opportunity to save it first. Then, a dialog will ask for the file to open.



This dialog works in the standard way. Select a file and click *Open* to open it, or navigate to another directory using the various means the dialog provides. It is also possible to specify the fully qualified pathname or a URL of a PNML file in the File name selector.

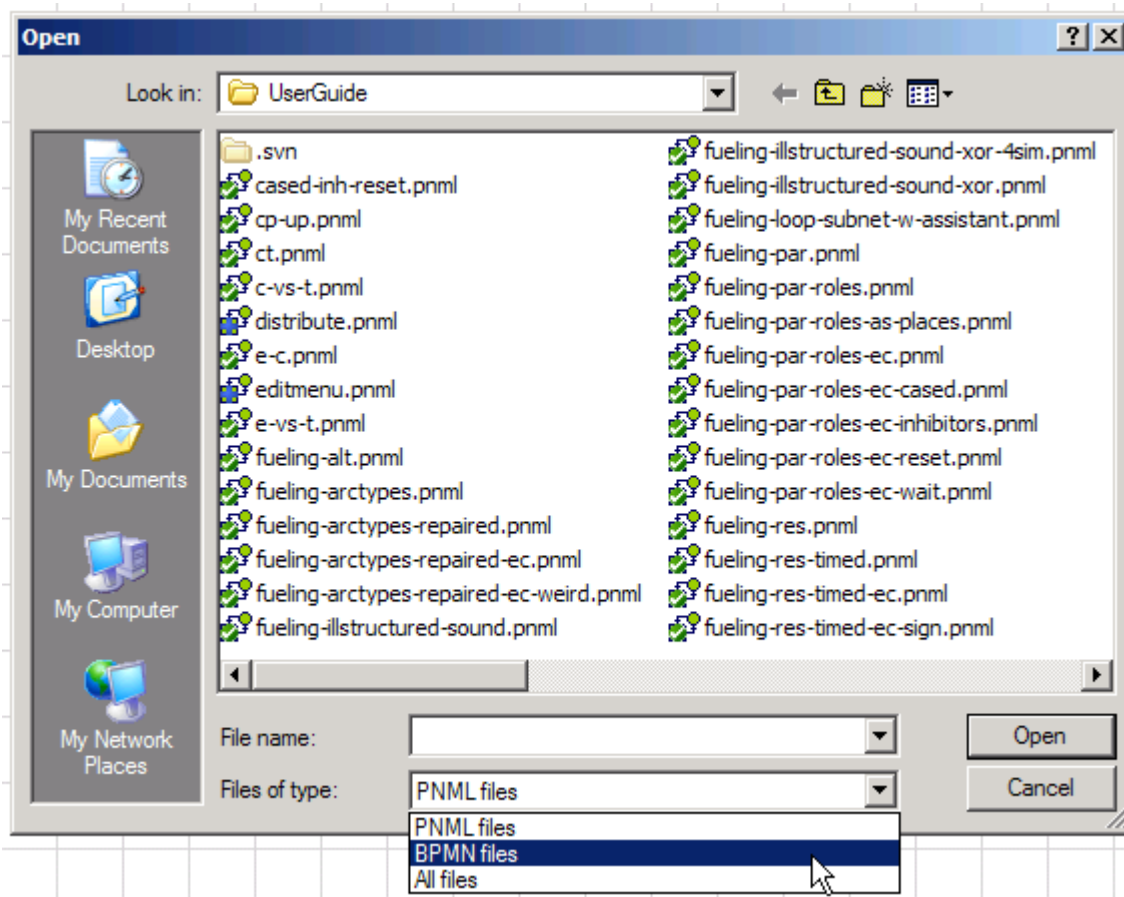
Alternatively, drag the file from Windows Explorer onto Yasper's editing surface.

Yasper saves its models in an XML-based file format conforming to the PNML standard²¹, with the extension *.pnml*. It can obviously read every model produced with Yasper, but its use of PNML also provides limited interchangeability with other process modeling tools. In principle, all basic nets (section 1.1) and some additional features can be interchanged with any other tool that can read or write PNML. In practice, interchangeability is still limited, since tools disagree on how to interpret or employ the PNML standard. Therefore, Yasper may refuse to read or misinterpret PNML produced by other tools, while other tools may do the same on PNML produced with Yasper.

²¹ For the PNML standard, see <http://www.informatik.hu-berlin.de/top/pnml/>;
for Yasper's use of it, see <http://www.petriweb.org/specs/>.

2.15.2 Importing BPMN models

The *Open* dialog can also be used to import models in BPMN format. This format is used by a tool called Deloitte Industry Prints Process and Repository to create process models in the BPMN notation.²²




To import a BPMN model in this format, use the *Open* function, but in the *Files of type:* selector, select *BPMN files*. Instead of PNML files, the list of BPMN files in the same directory will be shown. Pick one and click the *Open* button, or navigate to another directory with the methods provided by the dialog.

2.15.3 Exporting models to Microsoft Visio

To export a model to Microsoft Visio format, select *Export as Visio VDX ...* on the File menu.

²² For the BPMN standard, see <http://www.bpmn.org/>



The dialog offers two options:

- export to Visio 2002 VDX format
- export to Visio 2003 VDX format

Although the extensions are the same, the two document formats are incompatible. Visio 2002 and 2003 can partly read each other's VDX documents, but in doing so, throw warnings and lose information. Therefore it is important to select the correct format when exporting.

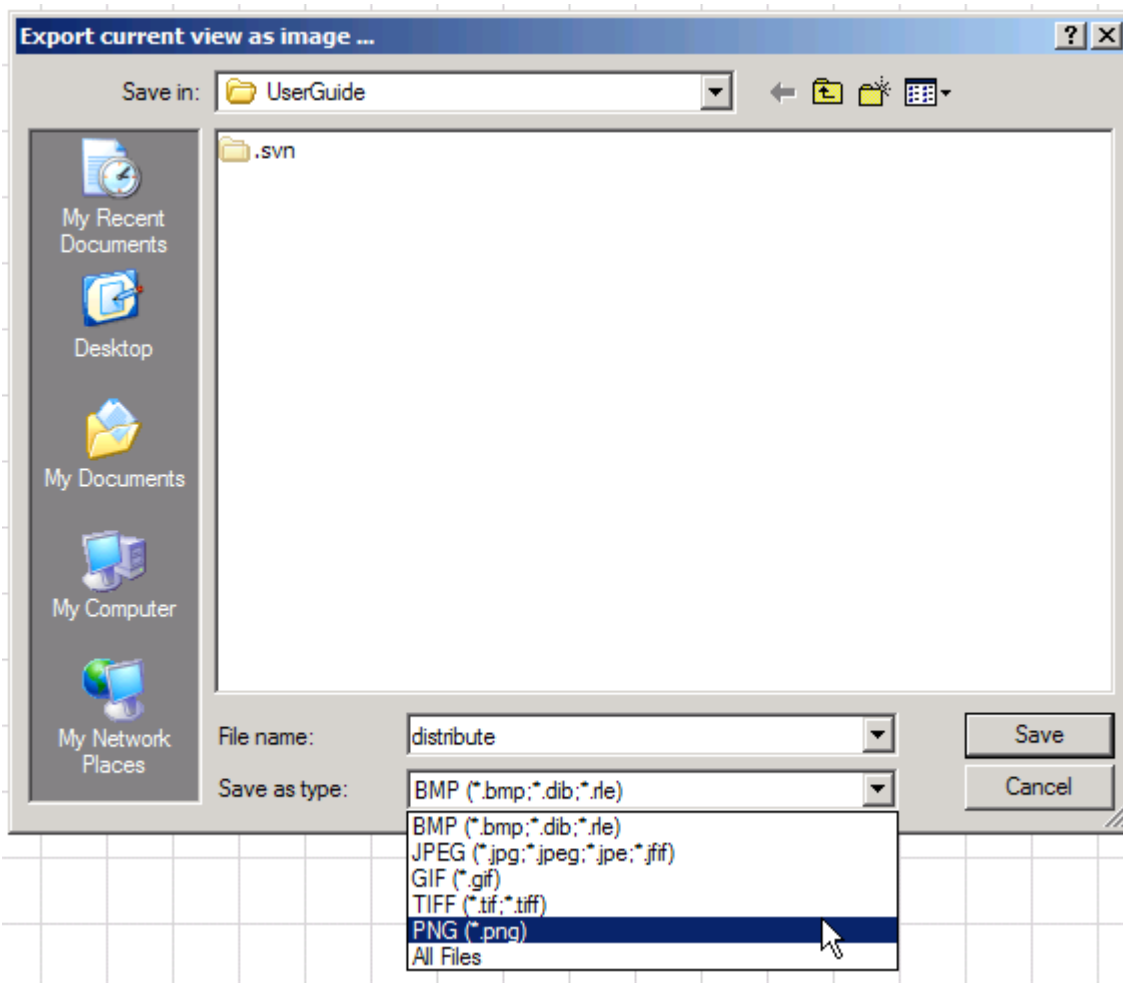
The translation provides a fairly accurate Visio representation of the Yasper model. If the document has subnets, they will be on separate Visio document pages, and the page symbol can be double-clicked to go to the page content.

The resulting Visio diagram is not intended as a replacement of the Yasper model:

- some information, e.g. on roles, is not included;
- there is no support for making sure that subsequent edits of the Visio document will leave it a valid representation of a Yasper model, and
- there is no support for importing Visio diagrams back into Yasper.

2.15.4 Exporting diagram images

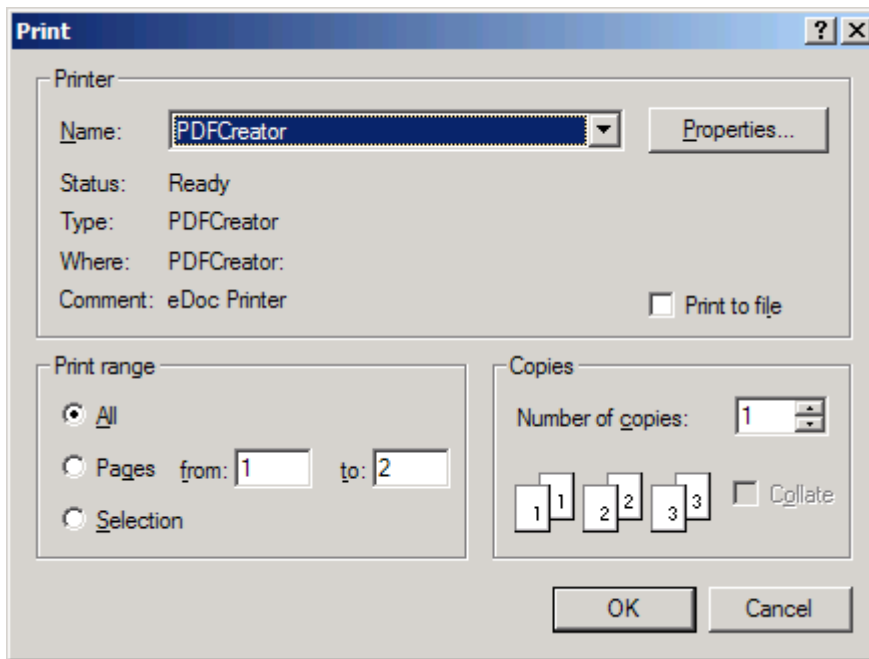
To export an image of the current (sub)net to file, select *Export as image ...* in the File menu. Various bitmap image formats are supported.



The image exported is the same as that copied by *Copy Image* (see section 2.14.2), except that no elements are selected.

2.15.5 Printing diagrams

To print a Yasper model, select *Print* in the *File* menu, click the *Print* icon in the toolbar, or press Ctrl-P. The print dialog will appear.



The images are identical to those produced by *Export as image* (see previous section), except that they are automatically scaled to fill the printed page.

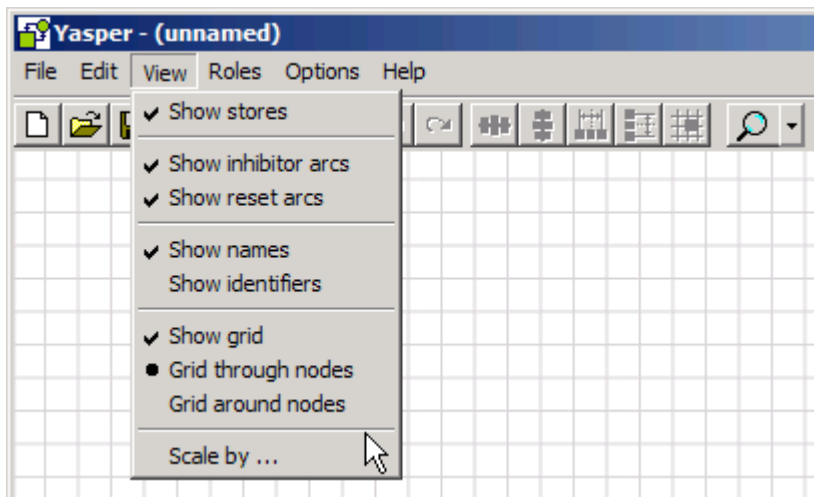
The *Print range* determines what to print:

- select *All* to print the main net and all subnets, each on a separate page;
- specify a page range to limit which of the nets are printed;
- select *Selection* to print the net that a *Copy* operation would copy (see section 2.14.2).

For a closer look at what is going to be printed, use *Print preview* on the *File* menu (*Shift-Ctrl-P*).

2.16 View menu options

Use the *View* menu to configure how the model is displayed.



Settings are remembered; future Yasper sessions will use them.²³

- When *Show stores* is off, data stores (cf. section 1.8) are omitted from the diagram. (They remain present in the model.)
- When *Show inhibitor arcs* is off, inhibitor arcs (cf. section 1.5.2) are omitted from the diagram. (They remain present in the model.)
- When *Show reset arcs* is off, reset arcs (cf. section 1.5.3) are omitted from the diagram. (They remain present in the model.)
- *Show names* displays the user-assigned names on elements. Note that such a name does not always exist and does not have to be unique.
- *Show identifiers* displays the system-assigned identifiers on elements. An identifier always exists and is always unique. Like *Show names*, this setting not only affects the diagram, but also the appearance of elements in dialogs and reports; it is convenient to set this when explicit names have not yet been assigned.
- *Show grid* displays a grid on the diagram surface; this is mainly useful for horizontal and vertical alignment, e.g. with the *snap to grid* feature (cf. section 2.18)
- *Grid through nodes* makes *snap to grid* snap elements to grid crossings;
- *Grid around nodes* makes *snap to grid* snap elements to grid tile centers.²⁴
- *Scale by ...* allows the scale factor of the diagram to be set (see section **Error! Reference source not found.**).

Most settings only affect the appearance of the diagram; the *Show names* and *Show identifiers* settings also affect how elements appear in the various dialogs used during editing, and in the

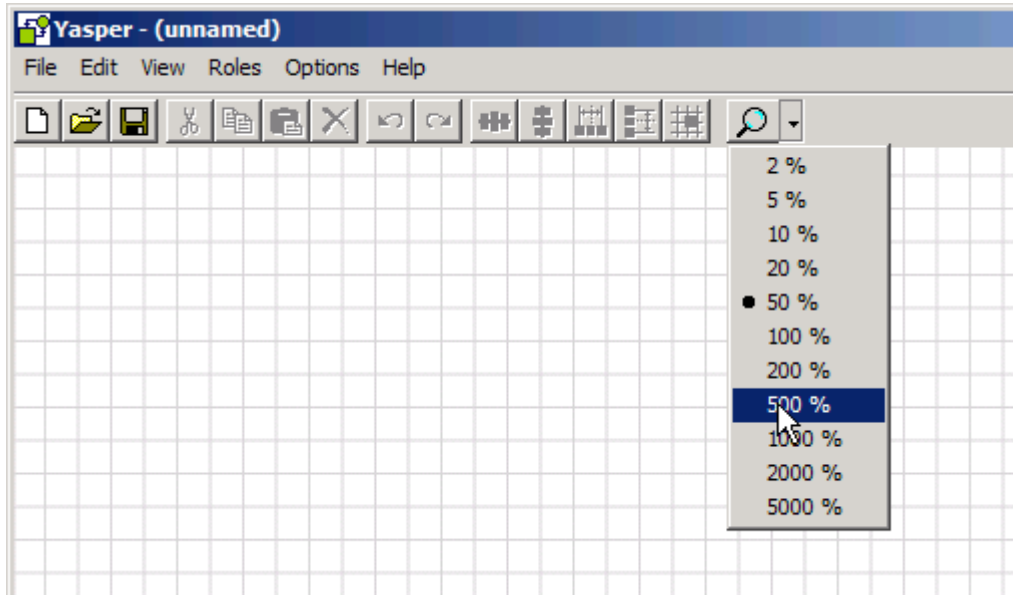
²³ However, if multiple Yasper sessions are running at the same time, changing a setting in one will not affect the other.

²⁴ *Grid around nodes* is the layout style used for most board games; Go players will prefer the *Grid through nodes* layout.

automatic simulation report. Names are assigned by the user, optional, and not always unique, while identifiers are machine-assigned, required and unique.

2.17 Zooming in and out

The diagram can be displayed at different scale factors. To change the scale factor, choose the *Scale by ...* item in the *View* menu, or click on the magnifying glass icon in the toolbar. Then pick the scale factor to use.

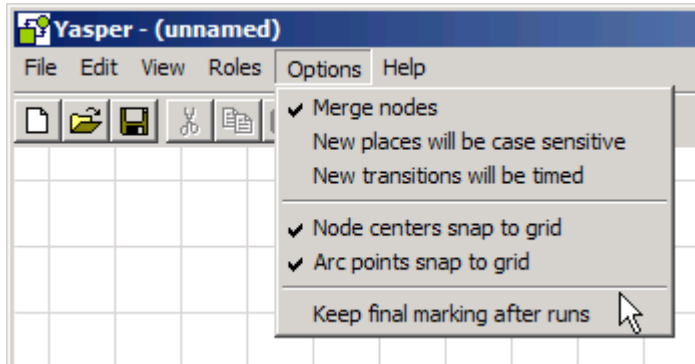


Like the other *View* settings, the scale factor is a global setting that is saved as a user preference. It does not depend on the model being shown and is not saved with the model.

Zooming in and out keeps the same point in the diagram in the center of the view. In the next version of Yasper there will be support for zooming in on a specific point.

2.18 Options menu options

The user preferences in the *Options* menu affect the way editor behaves. Like *View* settings, they are remembered for future Yasper sessions.



- *Merge nodes* makes a node melt away into a compatible node when moved onto it; for details see section 2.4.
- *New places will be case sensitive* determines whether places dragged from the building blocks (see section 2.1) are case sensitive.
- *New transitions will be timed* determines whether transitions dragged from the building blocks will be timed.
- *Node centers snap to grid* determines whether after moving selected elements, node centers will automatically snap to the gridline crossings (see section 2.3).
- *Arc points snap to grid* determines whether after moving selected arcs, their support points (see section 2.7) will snap to the grid lines.
- *Keep final marking after runs* determines what happens when switching from manual or automatic run mode back to editing mode. If this setting is disabled, editing will continue with the net as it was before simulation started; if it is enabled, it will continue with the token marking that was in effect in simulation mode.²⁵

²⁵ Except that if tokens were being processed by transitions, they will be forcibly produced, since tokens cannot reside in transitions during editing.

3 Manual simulation

3.1 What is simulation?

As explained in section 1.1.2, we know exactly what it means for a process described in Yasper to execute, since it is precisely defined (in section 1.6.4.4) when a process execution step is possible, what happens when it starts, and what happens when it finishes. Therefore, at any time it is unambiguously determined what may happen.

It is not always unambiguously determined what *will* happen; many choices remain open:

- multiple transition executions may be possible at the same time;
- for timed transitions, processing time may vary;
- for transitions with roles, multiple roles may be available;
- for choice elements (XORs), multiple inputs and/or outputs may be available.

Therefore it can be very instructive to see execution happen in a process model. This is known as *simulation*, since it simulates execution of the real process being modeled.

Yasper offers two methods of simulation:

- manual run: the user picks most of the choices by hand
- automated run: the system picks all of the choices at random

This chapter describes the facilities for manual simulation; automated simulation is described in chapter 4.

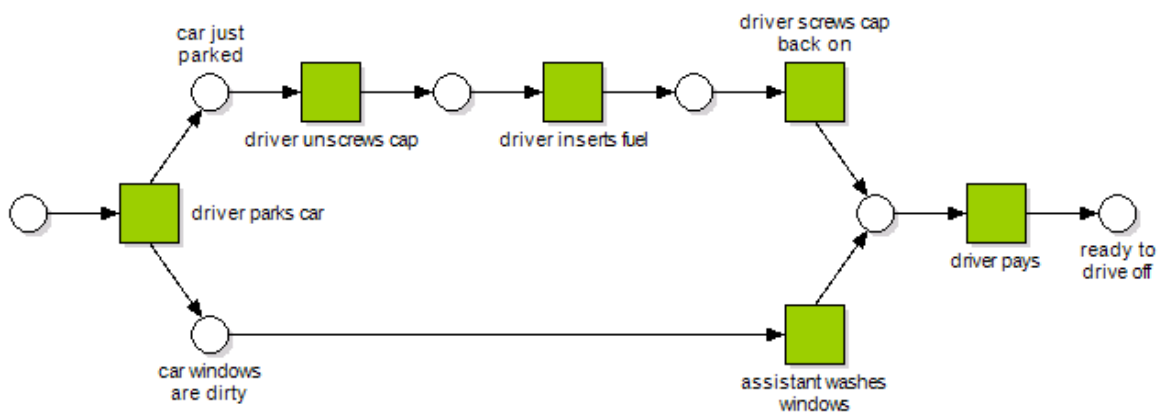
In manual simulation, the user executes the process step by step, making choices along the way. This is useful for showing that certain situations can be reached – for example, deadlocks. For quantitative analysis of a process model, use automatic simulation.

3.2 Preparing for simulation

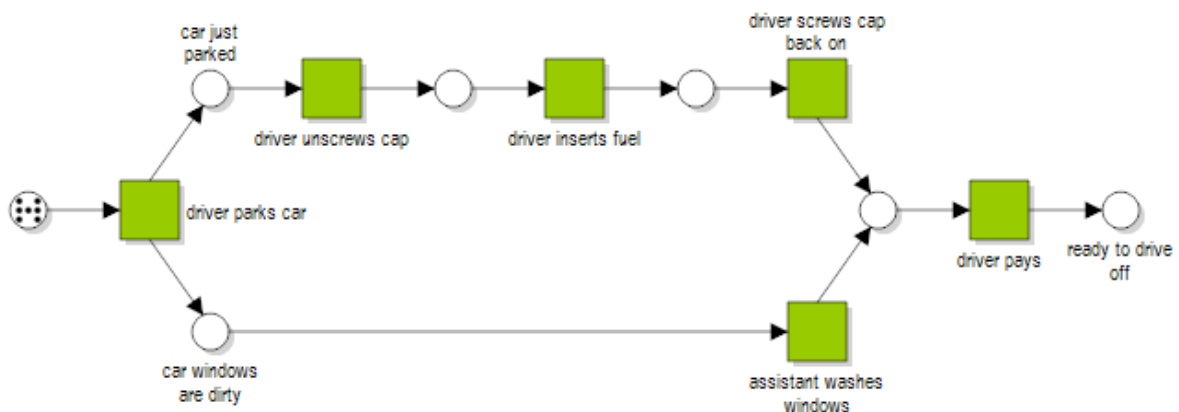
A model does not always represent the particular initial situation in which a run should start. For manual simulation, which does not employ time, such preparation mainly consists of choosing the right initial distribution of tokens. Models with initial places, such as those given in section 1.1 to 1.5, require that tokens be added to the initial places, or generated with additional transitions. Models with places that represent resources can be configured by setting various numbers of tokens in such places.

To perform such changes, go to *Edit* mode; in manual or automatic run simulation mode, the only way to modify the token marking is by executing transitions.

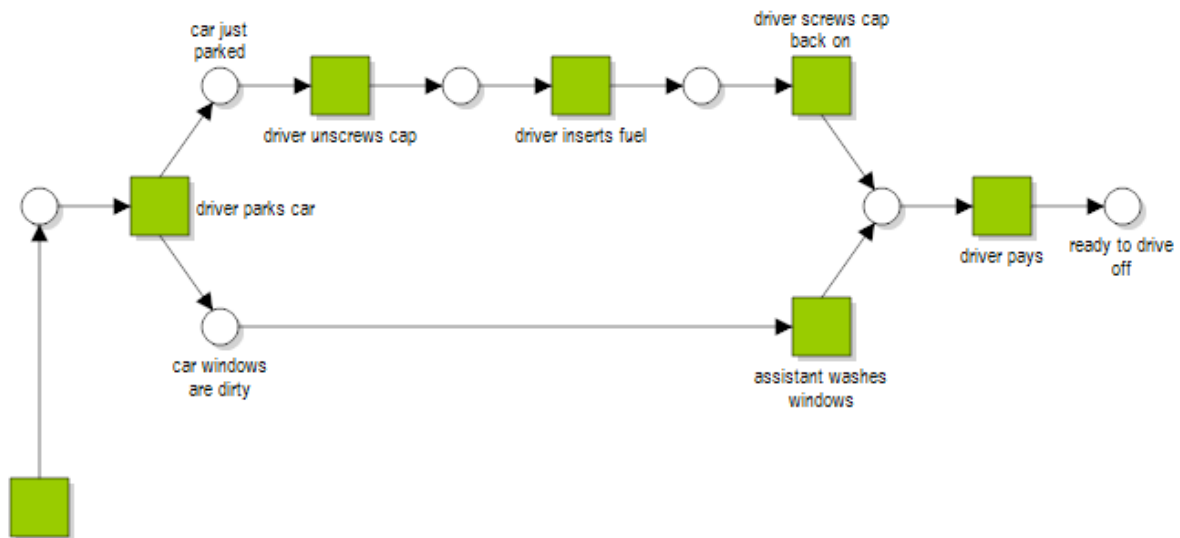
For instance, the following example from section 1.1.3.2 does not do anything when simulated:



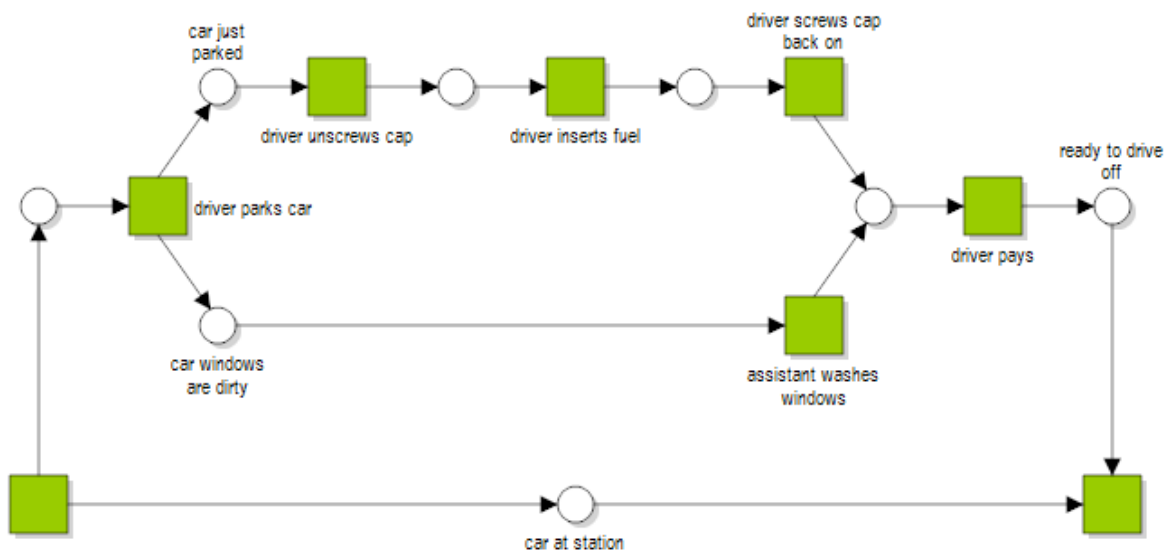
To allow simulation, we may add a couple of arriving cars:



or generate them with a new initial transition:

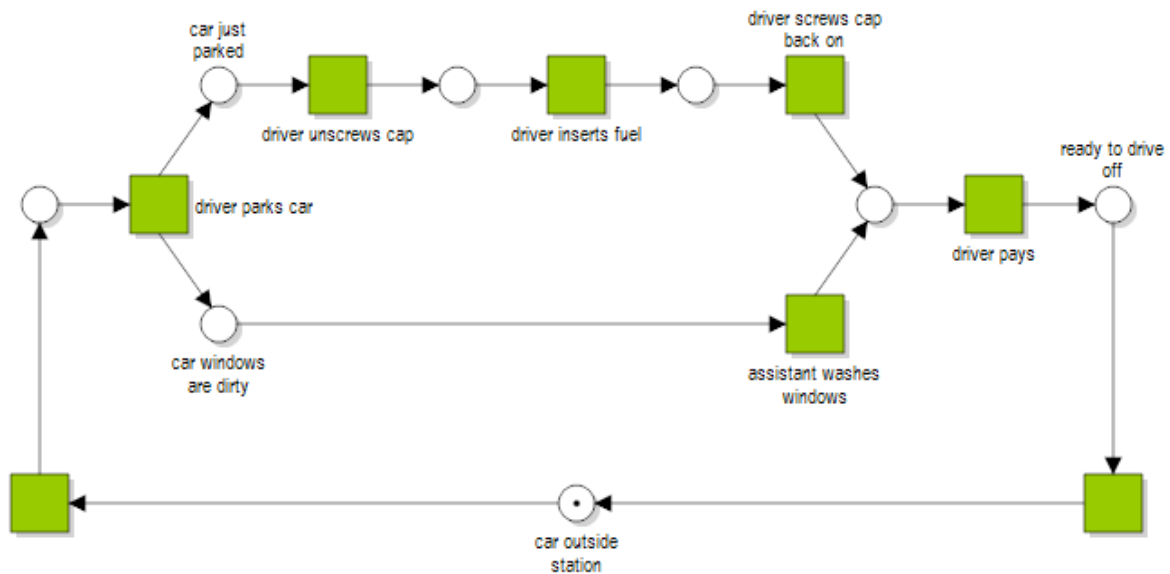


More additions can be made to allow more extensive testing:



If, during simulation, *ready to drive off* accumulates more tokens than *car at station*, we know something is wrong. (This will happen; the error in this model is explained in section 1.1.3.2.)

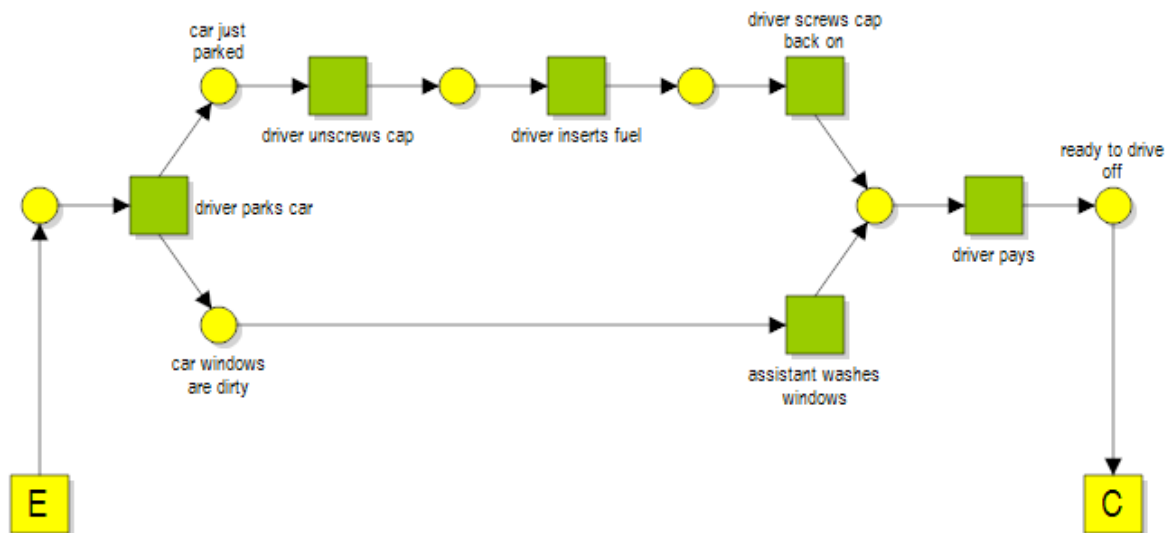
Another option is to add a test whether cars can pass through the station one by one:



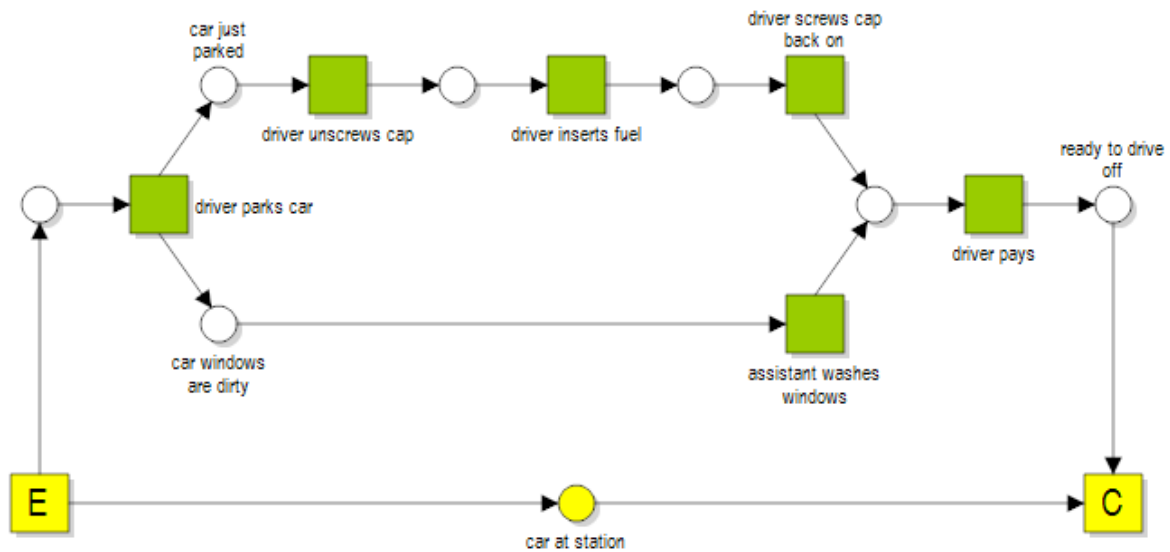
Here, automatic simulation will show an ever increasing number of cars flowing through the system, showing that something is wrong.

Many more sophisticated tests can be built by extending the model in more complex ways. Of course, care must be taken that the extensions meant to test the model do not interfere with its normal execution in unexpected ways.

Automatic simulation's reporting can also be useful, but it is only enabled when the net is a case sensitive workflow:



or contains one:

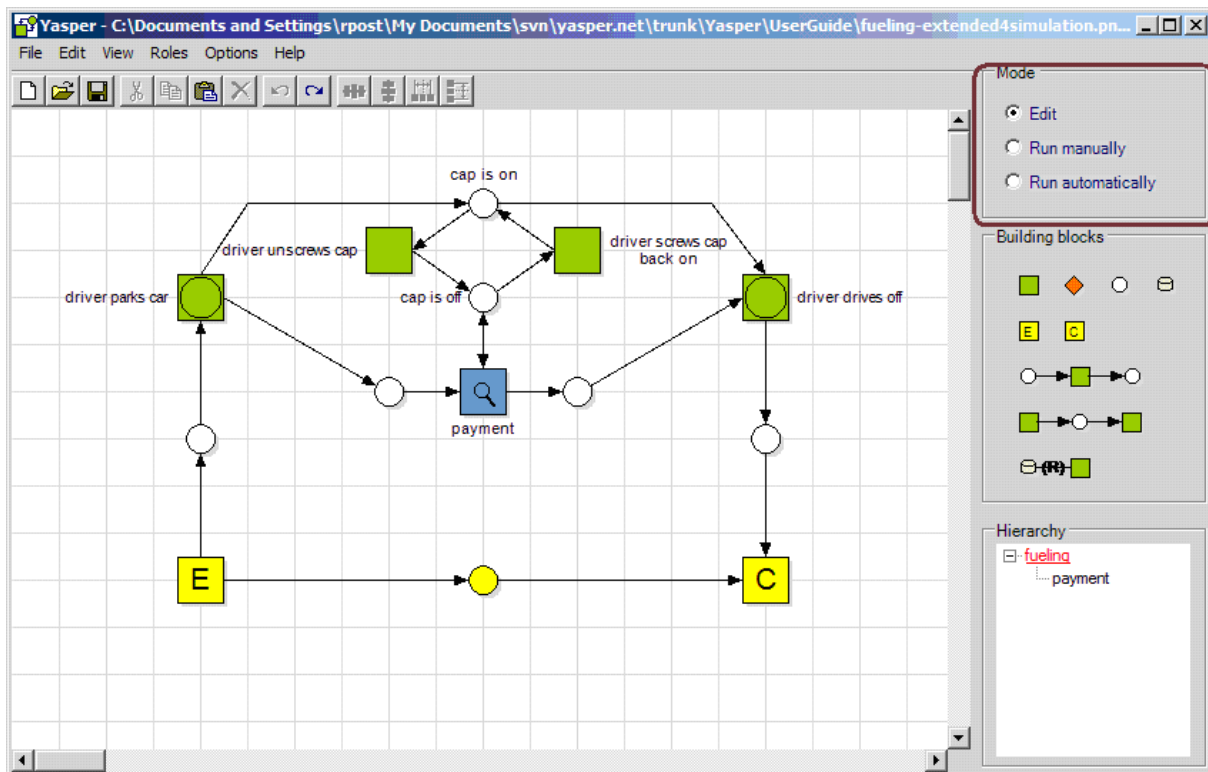


See section 1.6.4.4 for more examples.

Performance analysis can be added by equipping transitions with execution time, cost, and/or roles, as described in section 1.6.

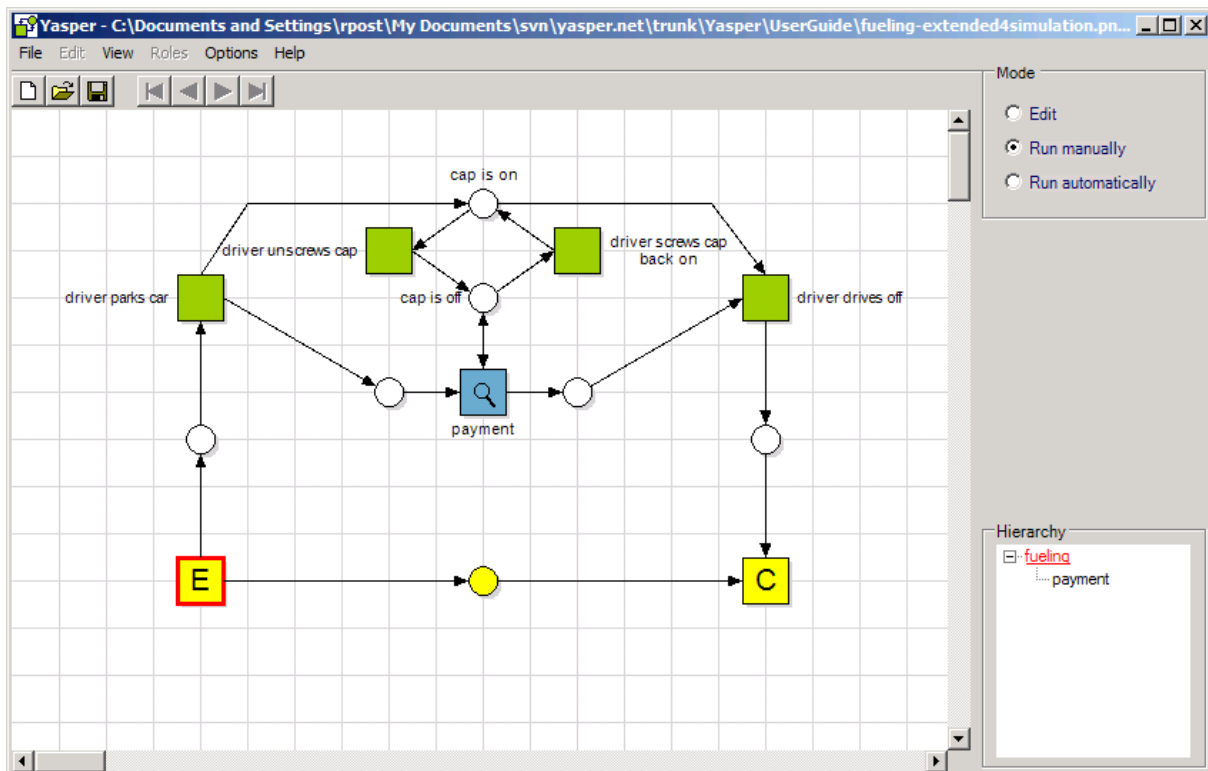
3.3 Starting manual simulation

In the right upper hand corner, three radio buttons control the mode Yasper is in.



To switch to manual simulation mode, click *Run manually*.²⁶

²⁶ This example describes the same process as the example in section 1.2; it was derived from it by adding an emitter-collector path, making some transitions timed, moving a part into a subnet, and changing the layout.



The display changes in various ways:

- there is no time in manual simulations, so timed transitions appear timeless;
- the elements that can execute are indicated by a red border;
- the toolbar gets four buttons, used in replaying the simulation.

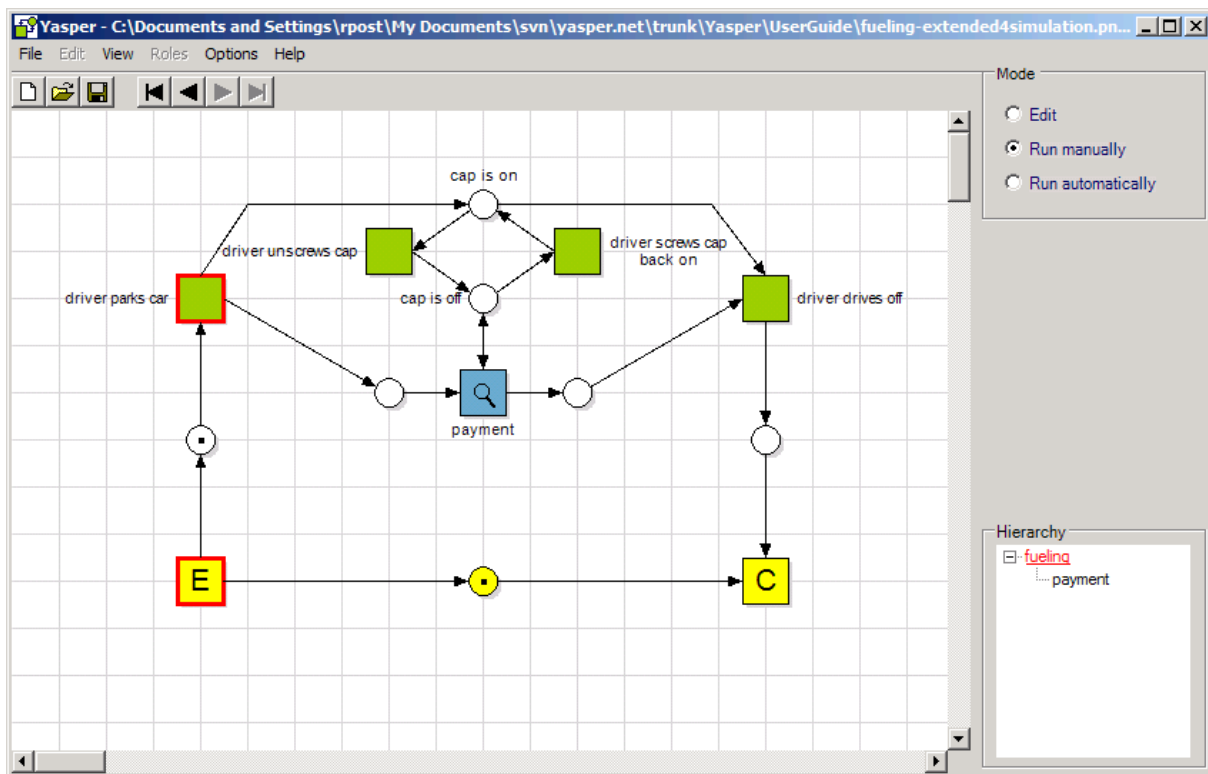
Editing operations become unavailable.

To switch back to edit mode, click *Edit mode*.

Some *File* operations can be invoked in manual simulation mode, but have the effect of switching to edit mode, since they do not make sense unless performed in edit mode.

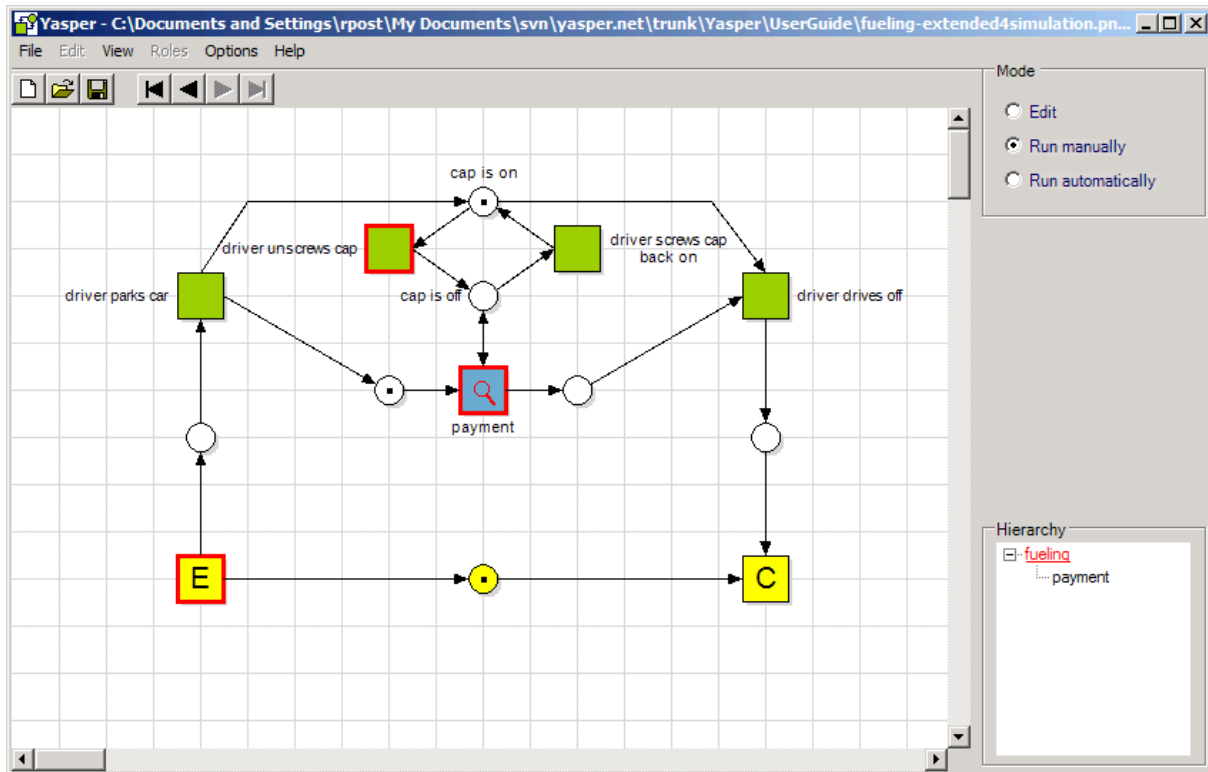
3.4 Executing a manual simulation

To execute a step, click on an element that is marked executable. In the above picture, that can only be the emitter; clicking it produces

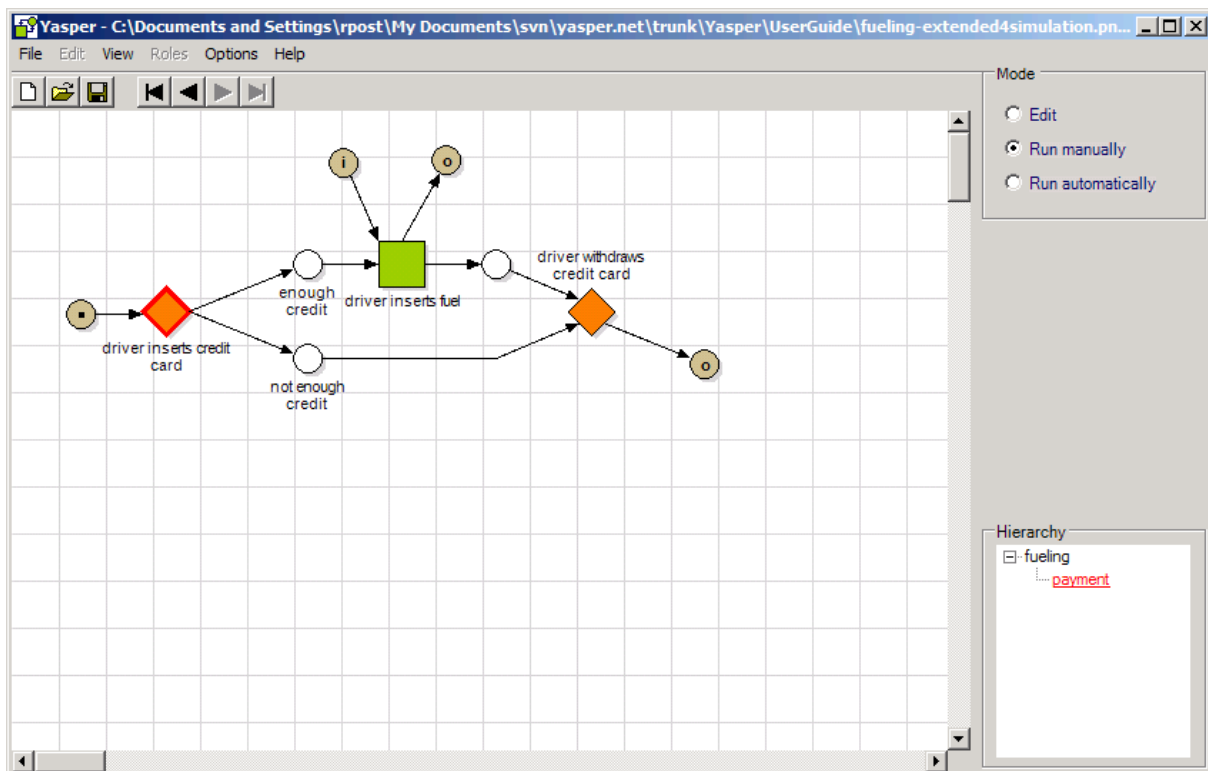


Repeat to continue. When multiple items are executable, select any of them.

Let's assume *driver parks car* is clicked.




A subnet is marked executable whenever any of its elements are. Click on the subnet to move inside.




To move up into the surrounding net, click on a pin.

Executing a XOR is usually a multi-step process.




To start executing one, click it.




If input can be taken from multiple places, they will be marked with red borders.

To pick one, click it.

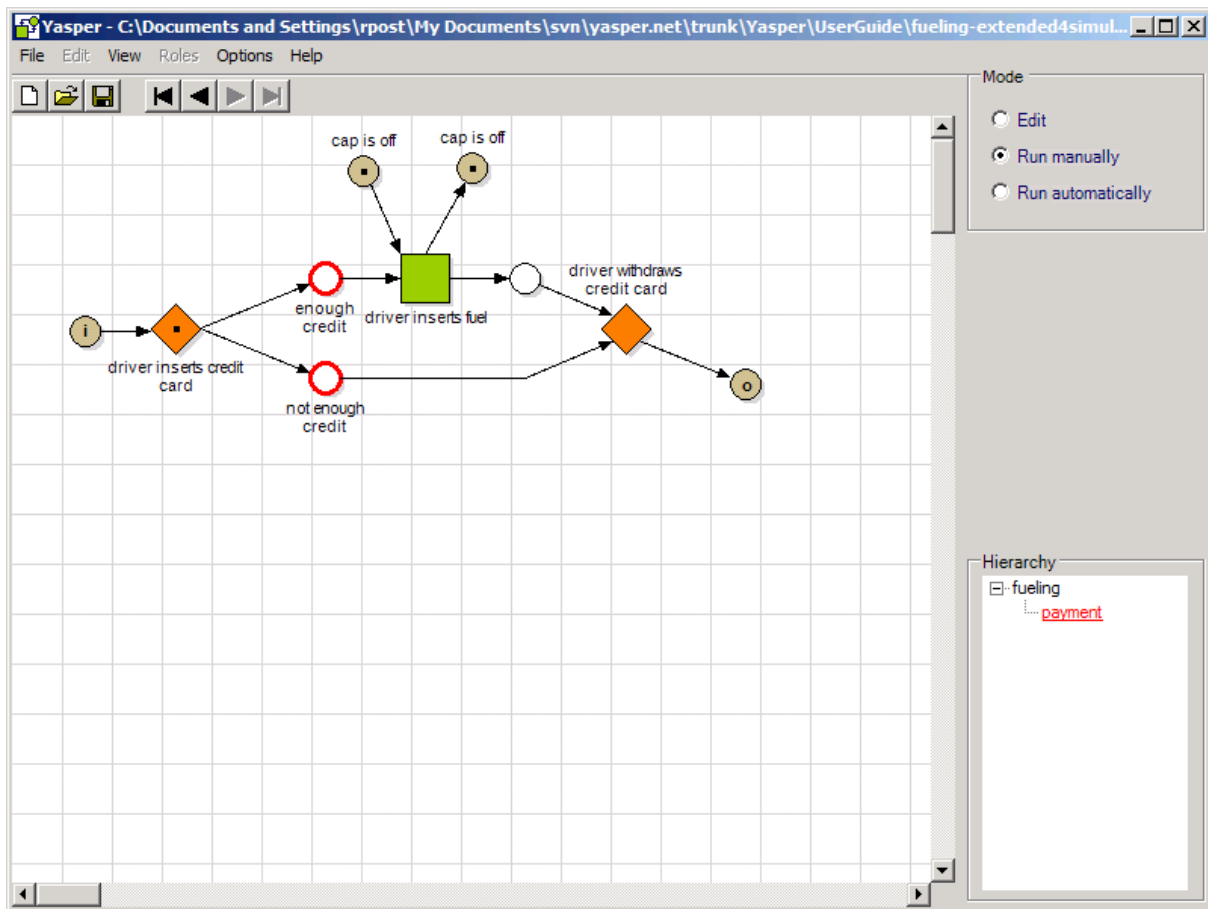


If output can be put into multiple places, they will be marked with red borders.

To pick one, click it.



If there aren't multiple inputs or outputs to choose from, the respective step is omitted. For instance, clicking the XOR in the example above, which only has one input, directly produces



3.5 Reviewing a simulation

Four controls in the toolbar appear in manual simulation mode:



Undo all steps – back to the start of the run



Undo the last step



Redo the last undone step

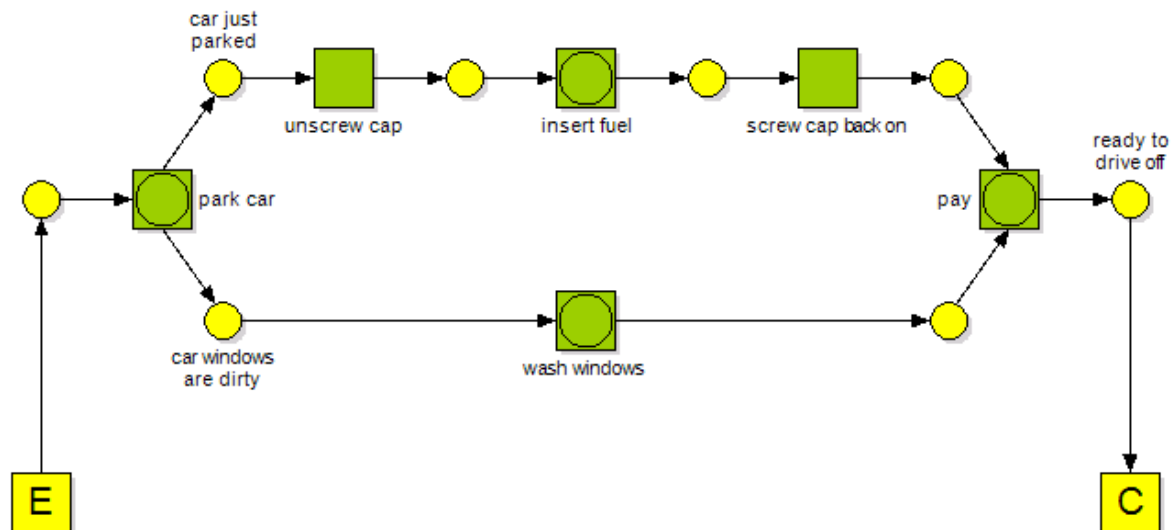


Redo all redoable steps – forward to after the last step taken

The buttons are grayed out when the function in question is not available. For example, right after taking a step, the first two are always enabled, since the step can be undone, while the other two are disabled, since there is nothing to redo.

3.6 Manual simulation with roles

When roles have been assigned to tasks (see section 2.12.2), manual simulation offers the choice which role to use when performing a task. Reconsider the first example of section 1.6.4.4, with roles assigned as in section 1.3:



Edit roles

roles	role	capacity
▶	driver	1
	assistant	2
*		

OK Cancel

Assign roles to tasks

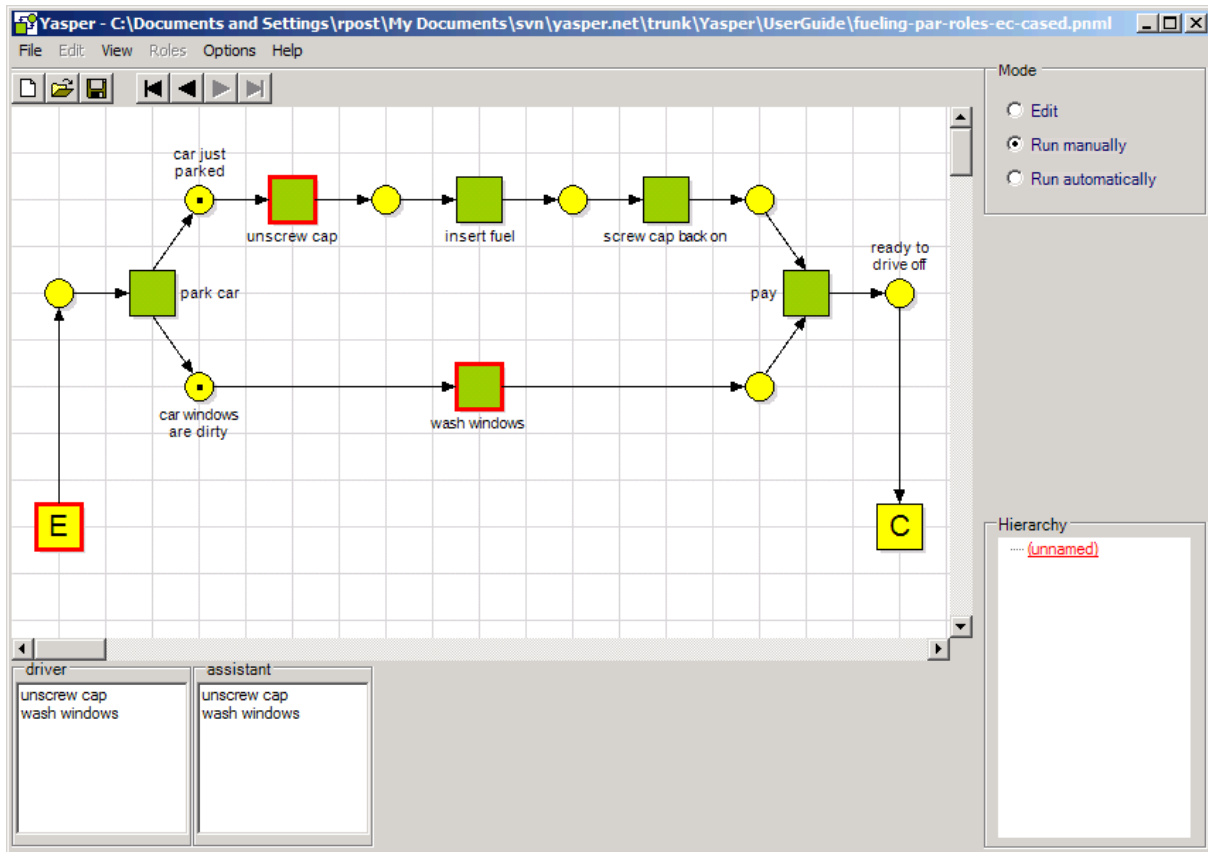
tasks shown:

☒ in whole net ☐ in current net

task	driver	assistant
unscrew cap	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
insert fuel	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
screw cap back on	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
wash windows	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
pay	<input checked="" type="checkbox"/>	<input type="checkbox"/>
park car	<input checked="" type="checkbox"/>	<input type="checkbox"/>

OK Cancel

In manual simulation mode, the interface displays a box for every role. At any time the box lists the transitions and XORs that can execute with that role:



To execute a task with any available role, click on it in the diagram.

To execute a task with a specific role, click on the task in the box for that role.

4 Automatic simulation

As discussed in section 3.1, Yasper provides the user with two methods to execute processes in a process model:

- manually (see chapter 3)
- automatically (this chapter)

In an automatic simulation, cases for the process are generated and executed automatically. Statistics on execution times and resource utilization are aggregated and summarized in a report.

While in manual runs, choices on how to continue are made by the user, during automatic runs all choices are made randomly, depending on configuration parameters set by the user beforehand.

Automatic simulation serves two main purposes. First, many logical modeling errors are caught when automatic simulation refuses to run or produces unexpected results. Second, once a correct process model has been constructed, simulation can be used to estimate performance and efficiency.

This chapter explains

- the general principles of automatic simulation (section 4.1)
- how to prepare for a run (sections 4.2 and 4.3)
- how to invoke and control a run (section 4.2)
- what the figures in the simulation report mean (section 4.5)
- switching views during simulation (section 4.6)

4.1 How automatic simulation works

Automatic simulation assumes *workflows*: process models that have the net effect of taking individual cases from a well-defined starting point to a well-defined end point. Before automatic simulation can be performed, the starting points, end points and intermediate points must be explicitly marked in the model with *emitters*, *collectors* and *case sensitive places*, respectively – section 1.6.4 provides full details and examples.

The simulator generates workflow cases in emitters, runs them through the process model, collecting statistics along the way, and summarizes them in a report. Some of the reported statistics pertain only to cases that have *completed* in collectors (see section 1.6.4.5).

Several user-configurable parameters influence the behavior of automatic simulation runs:

- *processing time* for transitions (see section 2.10.1 on how to set it)
- *output weights* on XORs (see section 2.10.2)
- *role capacities* and *role allocations* for either (see section 2.10.1)

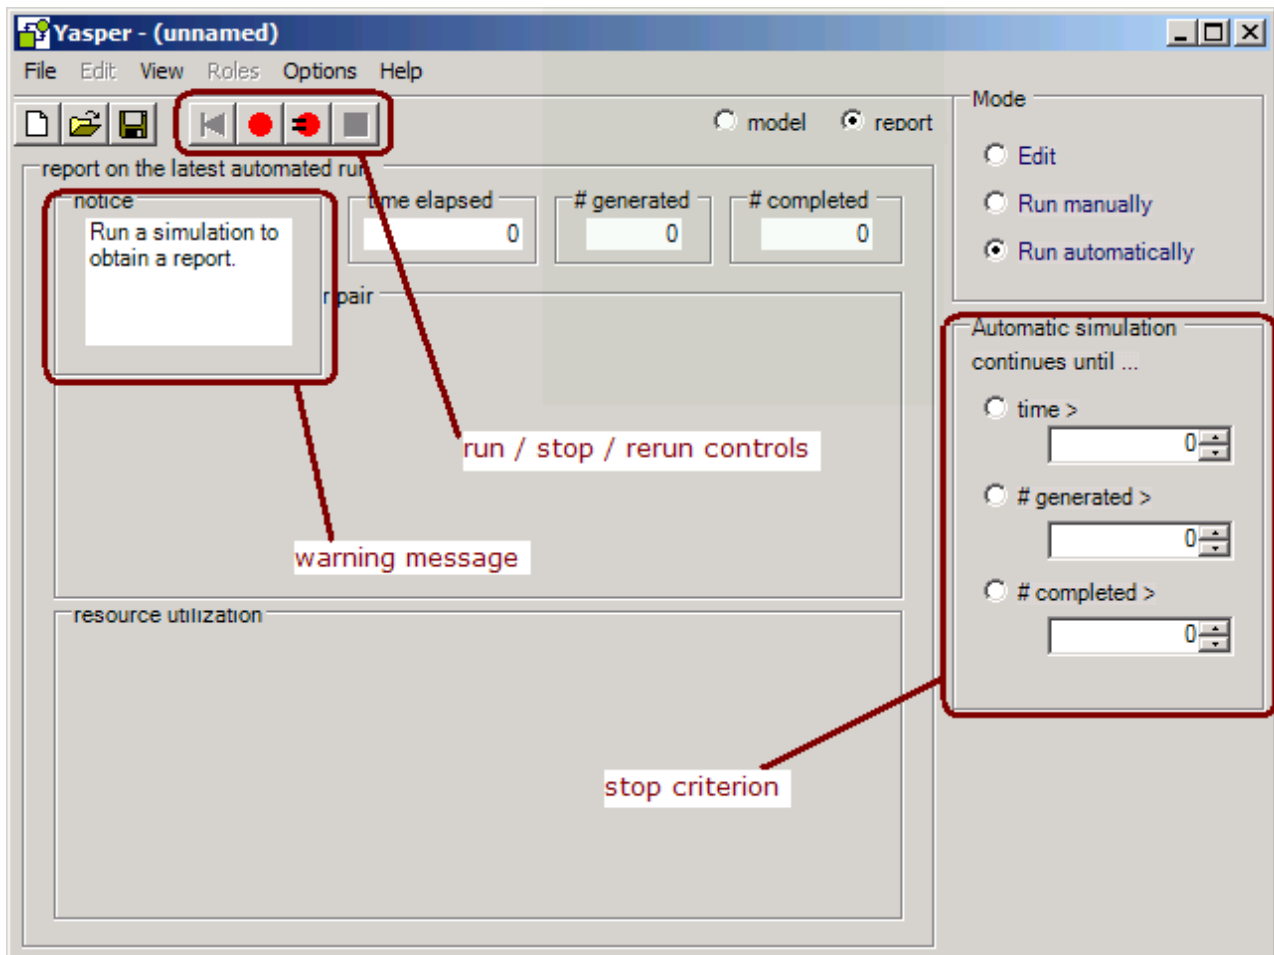
Another parameter merely helps summarize the behavior:

- *processing cost* for transitions (see section 2.10.1)

To modify these parameters, go to *Edit* mode; it is not possible to change them in run mode.

4.2 Preparing to run a simulation

Click on *Run automatically*. The controls for automatic simulation appear.



Three areas on the screen are of immediate interest:

- a *warning message* notifies whether simulation is possible
- *run controls* in the toolbar allow the simulation to be started
- a *stop criterion* to limit the simulation run

To get rid of the warning message, click on it.

To set a stop criterion, click the respective radio button and fill in a value:

- use *time >* to stop as soon as the given time is exceeded
- use *# generated >* to stop when the given number of cases has been generated
- use *# completed >* to stop when the given number of cases has been completed²⁷

²⁷ *Completed* in the sense of section 1.6.4.5, i.e. its last case token was collected by a collector.

4.3 Not being able to run a simulation

If simulation is impossible, all run controls are disabled.

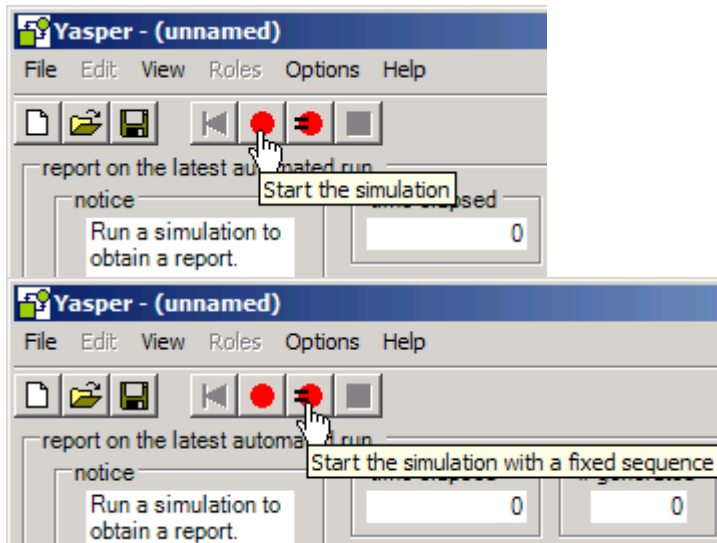
This will happen when

- no path of flow arcs leads from an emitter to a collector (so cases cannot complete)
- there is a timeless transition or XOR without inputs (so time cannot proceed)

To remedy this, click on *Edit* and change the model to remove these conditions.

For instance, basic Petri nets without case can often be prepared for simulation by adding an emitter, case-sensitive place and collector, and connecting them to the existing net.


4.4 Starting, halting and restarting a simulation



To start a run, click *Start the simulation* or *Start the simulation with a fixed sequence*. In the latter case, runs always take the same course when repeated.

These buttons are disabled when simulation cannot start. This is the case when


- a run is already in progress
- a halting criterion has been specified, and it is met
- no (first or next) step is possible (the simulation is in *deadlock*)
- simulation is impossible (as specified in the previous section)



During a run, the current time, number of generated cases, and number of completed cases are displayed in the respective fields.²⁸

To halt a run, click *Stop/Pause the simulation*. A report is generated.

²⁸ Whether or not these numbers always increase during a run depends on the model being simulated.



To continue a halted run, click *Resume the simulation*.

To revert to the initial situation, click *Back to start*.

4.5 The simulation report

When a simulation run halts, a report is generated.

report on the latest automated run

☒ show case list

time elapsed: 1013.39 # generated: 745 # completed: 443

report per emitter-collector pair

from	to	collected	completed	wait time	cycle time	work time	cost
(1) Car Damage Reported	(1) Car Damage Handled	196	196	124.63	134.1	9.46	283.89
(2) Car Damage Reported	(2) Car Damage Handled	247	247	113.48	117.86	4.41	64.88

resource utilization

rolename	% busy
Clerk	83.46
Expert	94.93
Expert2	94.27
Expert3	26.38

case list

case	from	to	completed?	wait time	cycle time	work time	cost
747	(1) Car Dama		no	1013.39	0	0	0
746	(2) Car Dama		no	1013.39	0	0	0
745	(1) Car Dama	(1) Car Dama	no	0.16	0	0	0
744	(1) Car Dama	(1) Car Dama	no	2.05	0	0	0

Automatic simulation continues until ...

☐ time > 0

☐ # generated = 0

☐ # completed = 0

Start Stop Reset Exit

A report can have three parts:

- *report per emitter-collector pair*
(always included, but empty when no case was completed)
- *resource utilization*
(only displayed when the model specifies roles)
- *case list*
(only displayed when *show case list* is checked)

The *report per emitter-collector pair* displays statistics on completed²⁹ cases, aggregated by their emitter and completing collector. The columns display the following:

- *from* = the emitter's name and/or id³⁰
- *to* = the completing collector's name and/or id
- *collected* = the total number of case tokens of cases with this emitter that ever arrived at this collector, including tokens for cases that haven't completed, or have completed in a different collector
- *completed* = the number of cases with this emitter completed by this collector

²⁹ *Completed* in the sense of section 1.6.4.5, i.e. their last case token was collected by a collector.

³⁰ Whether name and/or id are shown depends on the *Show names* and *Show identifiers* options in the View menu, cf. section 2.16.

All further quantities are averages on the completed cases only:

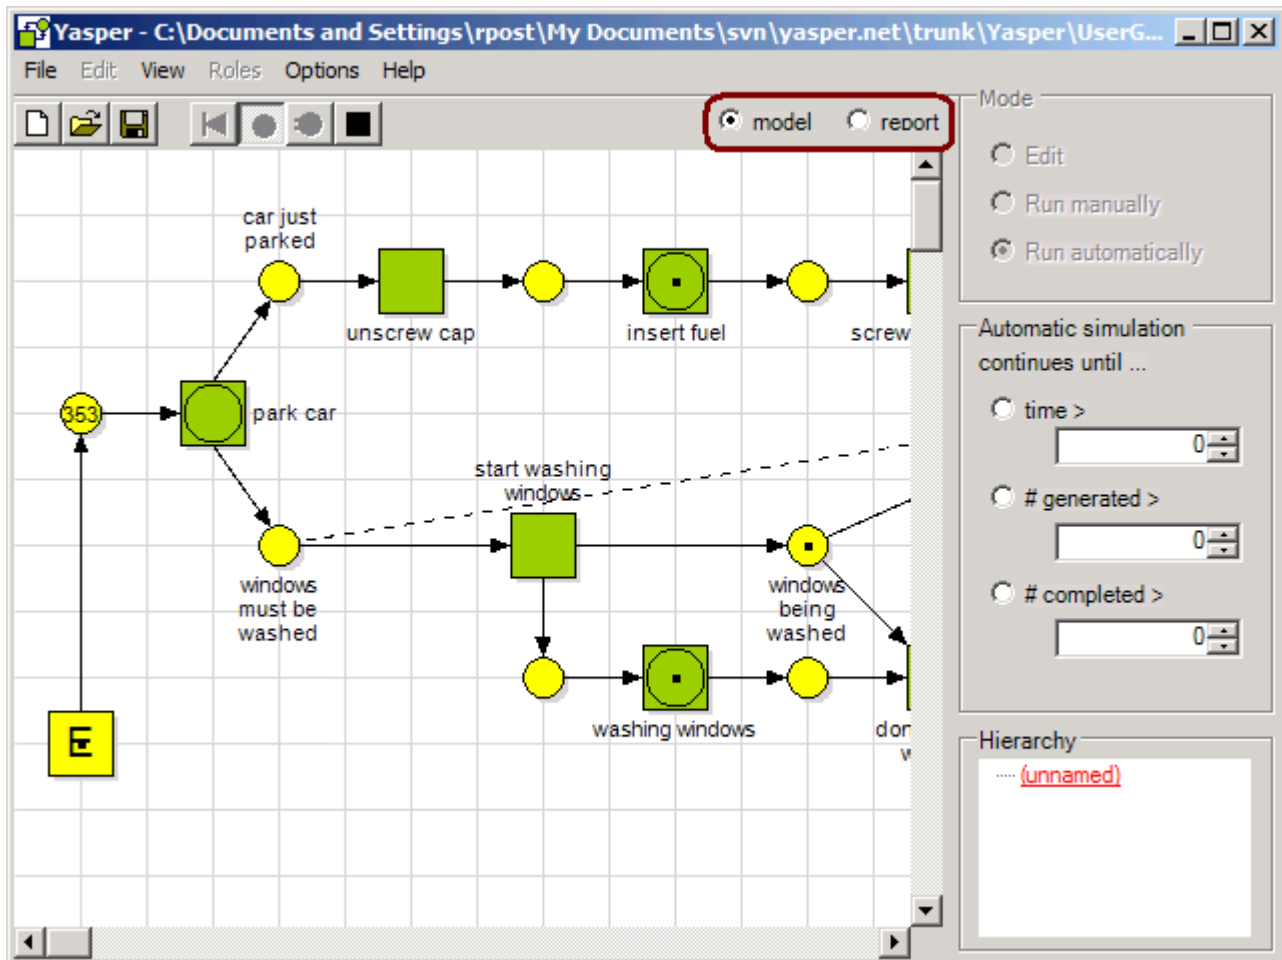
- *work time* = the average amount of processing time spent on a complete case by tasks; if parallel processing was done on the case, the processing times of all case tokens are added together
- *wait time* = the average amount of time a case was kept waiting, i.e. during which it had one or more tokens and no transition was processing any of them
- *cycle time* = the average amount of time a case was in the system, measured from leaving its emitter to being completed at its collector
- *cost* = the average processing cost of a case

The *resource utilization* report displays, for each role, what percentage of time its instances were involved in processing a task. All cases are considered, not only the ones that have completed.

The *case list* displays information on each individual case, whether completed or not. The meaning of the columns is the same as for the first table, except that it isn't aggregated.

4.6 Switching views and modes

To the right of the toolbar, two buttons appear in automatic run mode. These can be used to switch views between the reports and the diagram.



To see the token flow while a simulation is running, click on *model*.

To switch back to viewing the reports, click on *report*.

This can be done before, during, or after a run.

Switching views in automatic run mode is not the same thing as switching to *Edit* or *Run manually* and back. The former does not affect the simulation state, and can even be performed while a run is in progress. The latter halts the simulation, and either rewinds to the initial marking, or, when *Keep final marking after runs* is set in the *Options* menu, modifies the state to be a valid marking, by forcibly terminating any processing being done in transitions, since tokens cannot reside inside transitions in either mode.

Some *File* operations can be invoked in automatic simulation mode, but have the effect of switching to edit mode, since they do not make sense unless performed in edit mode.

5 Index

(This index is still incomplete.)

A

activity diagrams13
 alternative execution..10
 arc 7, 8
 in editing39
 arc types, special.....20
 arc weights
 output weights23
 arc, flow.....21
 arc, inhibitor
 inhibitor arc21
 arc, reset
 reset arc21
 arcs, combined
 combined arcs
 arc multiplicity 20
 automatic simulation,
 extensions for23

B

basic net.....7
 biflow arc.....20
 BPEL13
 BPMN14
 branching10, 13

C

capacity.....16
 case sensitive place....25
 choice10
 choice element:14
 complete properly.....28
 concurrent execution .10
 condition.....8
 correctness12
 cost per time unit24

D

deviation23

E

emitter26
 execution8
 execution rule8, 11
 explosion21

F

fixed cost.....24
 flow arc21
 flowcharts13, 14

H

hierarchy
 subnets18

I

incorrect model.....12
 input arc.....20
 interprocess
 communication.....13

L

loops14

M

marking8
 mean processing time 23
 merging nodes41
 mixed-case XOR26
 multiple tokens9

N

net 7
 node.....39
 merging41
 nondeterminism.....10

nondeterministic
 execution.....9

O

order of process steps 12
 output arc20

P

P/T net.....7
 parallel execution10
 Petri net7
 Petri nets6
 place7
 case sensitive25
 primary selection. 37, 41
 process charts14
 process execution ..8, 17
 process execution rule 14
 process execution with
 subnets19
 process flow logic11
 process model.....6
 processing cost24
 processing time23
 proper completion11
 proper completion.31

R

resource16
 resource contention ..11,
 13
 role16
 role capacity.....16
 role is not a place17

S

shortcut arcs13
 soundness11
 state machine.....13
 support point39

T

timed transition23
 token8
 token case25
 token game.....8
 transition7
 timed.....23

U

UML activity diagrams
 13, 14
 UML statecharts.....13

W

well-handled net.....11
 well-nested net11
 well-structured net11
 workflow.....11

X

XLANG13
 XOR.....14
 ambiguous26
 mixed-case26
 XOR split.....23

Y

Yasper website.....2
 Yasper's process models
 6